

Augmentation Robustness: Pass Rate across L0–L4 (CUDA→OpenMP, N=26 kernels/model, seed=42+L)

