

---

# ParBench: A Benchmark for Reliable Evaluation of LLM Parallel Code Translation

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Modern compute-intensive software is written against a rapidly changing ecosystem of accelerators, programming APIs, compiler stacks, and portability layers. As hardware and software platforms evolve, kernels in AI, scientific computing, simulation, graphics, and data-intensive workloads often need to migrate across CUDA, OpenMP, OpenCL, OpenMP target offload, and related parallel APIs. Large language models and autonomous coding agents are increasingly proposed as tools for this migration, but the field still lacks a reliable way to measure whether they can preserve the low-level parallel semantics that make such translations behaviorally valid under declared oracles, including thread indexing, synchronization, memory management, host-device coordination, and API-specific execution structure.

We present PARBENCH, a kernel-centric benchmark framework designed to isolate and measure LLM-based parallel API translation under executable, reproducible conditions. Unlike general code benchmarks that are largely sequential, parallel-code benchmarks that emphasize generation from natural-language prompts, or repository-level studies that confound translation with build-system reconstruction, PARBENCH fixes the surrounding build, run, and verification infrastructure through declarative benchmark specifications and asks models to translate only the computational kernels. The benchmark draws on multiple open-source HPC suites to cover representative cross-API translation directions among CUDA, OpenMP, OpenCL, and OpenMP target offload. To probe whether success reflects robust translation rather than surface-form memorization, PARBENCH also includes AST-driven intended behavior-preserving, baseline-validated source augmentation. We further evaluate PARBENCH on state-of-the-art open and proprietary LLMs, showing that it exposes persistent barriers to reliable parallel code translation, including direction asymmetry, multi-file coordination, incomplete API adaptation, and uneven robustness to source-level perturbations.

## 1 Introduction

Modern compute-intensive software is written against a rapidly changing ecosystem of accelerators, programming APIs, compiler stacks, and portability layers. Kernels in AI systems, scientific computing, simulation, graphics, and data-intensive workloads often need to move across CUDA, OpenMP, OpenCL, OpenMP target offload, and related parallel APIs as hardware platforms evolve or deployment requirements change. Large language models (LLMs) and autonomous coding agents are increasingly being explored for this migration task, but parallel API translation is fundamentally different from ordinary code translation. Correctness depends not only on preserving program intent, but also on maintaining thread-index arithmetic, synchronization semantics, memory-management patterns, host-device coordination, and API-specific execution structure—properties that go far beyond surface-level token replacement.




Figure 1: **PARBENCH isolates parallel API translation from repository-level confounds.** The benchmark pipeline starts by selecting representative parallel APIs and curating executable kernel specifications from multi-API HPC suites. Each specification fixes the build, run, and verification context, while the model is asked to translate only the computational kernel code. AST-driven source augmentations create surface-varied source variants for robustness testing, and translated outputs are evaluated through an end-to-end build–run–verify harness with conjunctive correctness checks. The final outcome taxonomy separates harness-passing translations (oracle PASS: compiled, ran, and satisfied the spec’s declared correctness oracle) from extraction, build, runtime, and verification failures, enabling both aggregate scoring and diagnostic analysis of where parallel translation breaks.

Can current models actually perform this translation reliably? Existing evaluations leave this question unresolved because they do not isolate parallel API translation as the measured capability. Widely adopted code-generation benchmarks such as HumanEval and SWE-bench Chen et al. [2021], Jimenez et al. [2024] are predominantly sequential and therefore do not exercise parallel programming semantics. Parallel code-generation benchmarks such as ParEval Nichols et al. [2024a] primarily evaluate synthesis from natural-language descriptions, rather than translation of existing implementations across APIs. Repository-level translation work exposes important integration challenges, but it also entangles kernel translation with build-system reconstruction, file organization, dependency management, and execution setup Davis et al. [2025], Wang et al. [2026]<sup>1</sup>. This leaves two core questions open: can models translate computational kernels across a broad set of parallel API directions when the surrounding infrastructure is fixed, and do successful translations reflect robust parallel-code translation rather than memorization of widely circulated benchmark code?

We present PARBENCH, a benchmark framework designed to answer these questions by isolating parallel kernel translation from repository reconstruction. When evaluating parallel code translation, build-system configuration, file scope, run arguments, and verification criteria vary across kernels and APIs; failures in any of these can confound assessment of translation quality itself. PARBENCH fixes the surrounding build, run, and verification infrastructure through declarative benchmark specifications, asks models to translate only computational kernel files, and evaluates each generated kernel with an end-to-end build–run–verify harness using conjunctive correctness checks. The corpus is survey-grounded: we screened 71 open-source HPC repositories spanning 29 parallel programming APIs and retained 35 that met inclusion criteria for kernel-level co-occurrence analysis (Appendix Figure 8). From these, we selected five benchmark suites that provide equivalent multi-API kernels and automated correctness checks (Section 3.1). To test whether successful translations reflect robust translation rather than surface-form memorization, PARBENCH includes an AST-driven augmentation engine that generates source variants intended to preserve behavior and validated against the baseline harness on the retained subset.

This paper makes three contributions:

- **Parallel code translation benchmark.** We introduce PARBENCH, a benchmark for LLM-based parallel code translation that provides executable kernel specifications spanning five HPC benchmark suites, six standard translation directions among CUDA, OpenMP, and

<sup>1</sup>An extended comparison with code-generation benchmarks, repository-level translation studies, HPC-specific LLM evaluation, and robustness benchmarks in Appendix A.

OpenCL, and four OpenMP-target case-study directions. By fixing the surrounding build infrastructure as context, PARBENCH enables direct build-run-verify evaluation of kernel translation without the build-system-generation confound identified by repository-level work.

- **AST-driven augmentation engine.** We develop an augmentation engine with six AST-driven source-level transformations applied at four intensity levels (L1–L4). The engine produces structurally varied source variants to test whether models preserve parallel semantics beyond surface similarity to widely circulated benchmark code.
- **Empirical characterization.** Across three models and 2,262 valid translation records, we report a multi-suite, multi-direction evaluation of LLM-based parallel code translation under stochastic sampling with pass@ $k$  analysis. Pass@1 ranges from 23.9% (Qwen 3.5) to 62.7% (GPT-5.4 and GPT-5.3-codex) on 142 unique no-augmentation (L0) tasks, with the code-specialized GPT-5.3-codex statistically indistinguishable from the general-purpose GPT-5.4. Key findings include build-stage API adaptation as the main bottleneck, strong direction and file-boundary effects, generally stable performance under augmentation with model-dependent variation at higher levels, and limited gains from repeated sampling (with  $n=3$  samples, most failures remain hard failures across all attempts).

## 2 PARBENCH Framework

To better evaluate LLM-based translation of parallel programming APIs, PARBENCH defines a controlled, kernel-centric setting that separates parallel kernel translation from repository-level challenges such as build-system reconstruction. The framework realizes this setting through four components: (1) declarative specifications, (2) an AST-driven augmentation engine, (3) an evaluation pipeline that orchestrates LLM translation and records diagnostic failure categories, and (4) a build-run verify harness. Together, these components enable direct measurement of parallel API translation capability under end-to-end executability constraints while supporting controlled robustness analysis. Figure 1 shows the overall architecture.

### 2.1 Declarative Specifications and Kernel-Centric Translation

Each benchmark kernel-API variant in PARBENCH is encoded as a JSON specification (spec) that serves as a declarative contract defining what a passing execution means for that kernel-API variant under the declared oracle. The spec encodes all information required to build, run, and verify a benchmark without manual intervention, enabling fully automated evaluation. The schema is formally defined in JSON Schema and enforced by an offline validator. A spec declares the kernel’s identity, specifies the target API, and includes provenance details to ensure reproducibility by including the pinned repository and commit. The spec partitions source files into `prompt_payload` (source files shown to the LLM), `support_files` (build infrastructure and shared headers; header content is included in the prompt as read-only context for interface compatibility), and `verification_only` (reference implementations never shown to the LLM). A `translation_targets` field identifies the subset of payload files the LLM must rewrite, enabling the kernel-centric translation described in Section 2.3.

**Build, run, and verification.** The build block specifies compiler commands, toolchain variables, and expected executable path. The run block defines input configurations with arguments and timeouts. The verification block declares an ordered list of verification strategies—exit-code checks, stdout pattern matching, quantitative comparison, and cryptographic hash verification—evaluated in conjunction; all must pass for an overall PASS verdict (Section 2.4).

**Baseline results.** The `baseline_results` block records the output from running the original implementation on a reference platform, providing reference values from which verification strategy parameters are calibrated. Section 3 describes the construction of 87 verified evaluation specifications, with an example BFS specification shown in Listing 1.

### 2.2 Augmentation Engine: Probing Robustness via Code Perturbations

Widely used suites such as Rodinia Che et al. [2009] are likely present in LLM training data, so success on unmodified source may reflect memorization. The augmentation engine generates surface-

varied source variants by applying six AST-level transforms backed by the `libclang` API: condition operand swapping (e.g.,  $a < b \rightarrow b > a$ ), arithmetic form conversion (e.g.,  $x += 1 \leftrightarrow x = x + 1$ ), scope-aware variable renaming, typedef expansion, pointer-to-array notation conversion (e.g.,  $*(arr + i) \rightarrow arr[i]$ ), and static function renaming. The transform set is designed to preserve parallel structure (thread indexing, synchronization, memory management) while changing surface-level code features most likely to be memorized verbatim. Transforms that would alter parallel semantics (e.g., loop reordering) are deliberately omitted.

Five augmentation levels (L0–L4) control transform density. L0 is unmodified source. L1 applies one randomly selected transform to one candidate site. L2–L4 apply increasing fractions of applicable transforms ( $f=0.33, 0.66, 1.0$ ), each to the same fraction of candidate sites. Baseline validation supports the intended behavior-preserving design on the retained subset: all 87 non-`KNOWN_FAIL` specs pass at L1–L2, and 80 of 87 pass at all levels through L4. The 7 high-intensity baseline failures occur exclusively in `omp_target` variants, where condition swapping is associated with numerically different device-code behavior under the offloading toolchain. Variable renaming is disabled for files containing OpenMP pragmas. These validation failures are treated as augmentation/toolchain-specific exclusions rather than evidence about model robustness for the affected `omp_target` baselines. Level definitions and per-transform details appear in Appendix Table 8 and Appendix E.2.

### 2.3 Evaluation Pipeline: Kernel-Centric Translation

The key design choice is kernel-centric translation: each spec partitions files into three roles. Fixed build infrastructure (Makefiles, build commands, verification logic) is never part of the translation task. Read-only context—source headers from `support_files` and target-side infrastructure (non-kernel payload files, shared headers)—is provided to the LLM for interface compatibility but is not part of the translation output. Translated source files (`translation_targets`) are what the LLM must rewrite; for full-program translations (e.g., `CUDA ↔ OpenMP`), these include host-side API calls, whereas for kernel-only translations (e.g., `OpenCL .cl` files), only the device kernel is translated. This separation isolates parallel API translation from repository reconstruction—the confound that drives ParEval-Repo’s 0% pass rate on applications above 133 SLoC Davis et al. [2025].

The pipeline constructs a structured prompt containing the source code, target API and output filenames, the build command, and read-only infrastructure context, with prompt anonymization (generic file labels, stripped kernel names and source comments) to reduce benchmark identity leakage. For cross-API translations, argument and verification asymmetry are handled automatically: full-program translations use the source spec’s run arguments and match against a union of source and target output patterns, while kernel-only translations use the target spec directly since the host runtime is unchanged. Each LLM response is parsed via a four-tier extraction strategy (Appendix B) that progressively relaxes from explicit filename annotations to fuzzy proximity matching and single-file elimination. If any expected file cannot be recovered, the task is classified as `EXTRACT_FAIL`. Source–target pairs are further classified during analysis by translation cardinality—`single_file`, `multi_to_single`, `single_to_multi`, or `multi_to_multi`—enabling stratified difficulty analysis. The full prompt template is in Appendix F.

### 2.4 Harness: Build, Run, Verify

The harness is a three-stage pipeline where failure at any stage short-circuits subsequent stages. The **build stage** resolves the spec’s working directory, substitutes platform-specific paths, and invokes compilation. The **run stage** executes the compiled binary with spec-defined arguments and timeout. The **verify stage** applies verification strategies in conjunction—all must pass for a `PASS` verdict. Five strategy types are supported: `exit_code`, `stdout_pattern`, `stdout_exclude_pattern`, `numeric_comparison`, and `file_hash`. Accordingly, `PASS` denotes successful build, execution, and satisfaction of the declared verification oracle, not a proof of full semantic equivalence.

Conjunctive verification is essential. For example, a CUDA translation of Gaussian elimination compiles, runs to completion with exit code zero, and produces partial stdout—but omits the expected timing summary, indicating a systematic translation defect that compile-only (TehraniJamsaz et al. [2024]) or exit-code-only verification would miss. Each evaluation trial receives one of four failure labels—`BUILD_FAIL`, `RUN_FAIL`, `VERIFY_FAIL`, or `EXTRACT_FAIL` (when no compilable

code can be parsed from the LLM response)—or PASS, enabling systematic failure-mode analysis (Section 5.4).

### 3 Benchmark Curation

The benchmark corpus was assembled through a four-stage systematic process: surveying open-source HPC benchmark repositories, quantifying kernel-level translation opportunities across parallel APIs, filtering candidate kernels through automated build-run-verify checks, and verifying each selected kernel through the complete pipeline on the evaluation platform. This curation populates the Source Code component of the PARBENCH architecture (Figure 1) with kernel files from five benchmark suites spanning four parallel APIs—CUDA, OpenMP, OpenCL, and OpenMP target offload—with the first three serving as standard translation directions and OMP-target as a case-study extension (Section 4).

#### 3.1 Suite and Kernel Selection

We screened 71 open-source HPC repositories spanning benchmark suites, mini-applications, proxy applications, libraries, and microbenchmarks across 29 parallel programming APIs, retaining 35 that met inclusion criteria for detailed kernel-level analysis (Appendix C.1).

A central finding is that **repository-level counting dramatically understates the available material for parallel translation evaluation**: although only 6 repositories contain both CUDA and OpenMP implementations, kernel-level analysis reveals 472 independent CUDA–OpenMP translation pairs—a  $79\times$  multiplier (Appendix Figure 8). This gap motivates PARBENCH’s individual kernel-centric evaluation design, rather than evaluating repository-level translation.

Five suites were selected based on their open-source availability, multi-API kernel equivalence, build-run-verify automation, self-checking correctness labels, and domain diversity; detailed criteria and exclusion rationale are provided in Appendix D.2–D.7. **Rodinia** Che et al. [2009] is the largest contributor, providing 60 specs across 22 kernels over CUDA, OpenMP, and OpenCL; 53 pass baseline verification and 7 are KNOWN\_FAIL. **HeCBench** Jin and Vetter [2023] provides the largest initial multi-API pool: from 522 kernels, a structured funnel yields 10 curated kernels—five with CUDA, OpenMP, and OMP-target variants, and five with CUDA and OMP-target only—totaling 25 specs. **XSbench** Tramm et al. [2014] and **RSBench** Tramm and Siegel [2014] contribute nuclear-physics proxy kernels, each with 4 specs, while **mixbench** Konstantinidis and Cotronis [2017] contributes 3 GPU roofline micro-benchmark specs. All XSbench, RSBench, and mixbench specs pass baseline verification.<sup>2</sup>

#### 3.2 Evaluation Corpus

The 87 verified-PASS specs constitute the evaluation corpus (Table 1). Kernel complexity spans more than an order of magnitude—80 to 3,304 SLoC, median 271—and 31 of 35 kernels (89%) exceed the 133-SLoC scale at which prior repository-level translation reports 0% pass@1 Davis et al. [2025]. Of the 96 curated specs (87 PASS plus 9 KNOWN\_FAIL), 25% require multi-file translation; CUDA has the highest multi-file rate (51%) owing to separate host and kernel sources, versus 12% for OpenMP and 13% for OpenCL (Appendix Table 7). These properties place the corpus well beyond single-function benchmarks while retaining fixed build–run–verify infrastructure. The corpus draws predominantly from Rodinia (53 of 87 non-KNOWN\_FAIL specs, 61%), so aggregate pass rates partly reflect its characteristic stencil, graph-traversal, and simulation-loop kernels.

## 4 Experimental Setup

**Models.** We evaluate three models: Qwen 3.5 397B-A17B, a 397B-parameter MoE open-weight model accessed via Together AI; GPT-5.4, a general-purpose proprietary model accessed via Azure OpenAI; and GPT-5.3-codex, a code-specialized model optimized via reinforcement learning on software engineering tasks OpenAI [2026], also accessed via Azure OpenAI.

<sup>2</sup>Corpus commits: Rodinia 9c10d3ea; HeCBench 22785cdd; XSbench ba08e522; RSBench 34b64478; mixbench 32edeca9.

Table 1: Evaluation corpus. SLoC is measured on the CUDA source variant when available. Multi-File denotes specs requiring more than one translated source file. KF = KNOWN\_FAIL.

| Suite        | Kernels   | Specs     | PASS      | KF       | APIs                    | SLoC Range      | Med.       | Multi-File         |
|--------------|-----------|-----------|-----------|----------|-------------------------|-----------------|------------|--------------------|
| Rodinia      | 22        | 60        | 53        | 7        | CUDA, OMP, OCL          | 195–3,304       | 334        | 21/60 (35%)        |
| HeCBench     | 10        | 25        | 23        | 2        | CUDA, OMP, OMP-tgt      | 80–235          | 180        | 0/25 (0%)          |
| XSbench      | 1         | 4         | 4         | 0        | CUDA, OMP, OCL, OMP-tgt | 1,390           | –          | 2/4 (50%)          |
| RSBench      | 1         | 4         | 4         | 0        | CUDA, OMP, OCL, OMP-tgt | 1,016           | –          | 1/4 (25%)          |
| mixbench     | 1         | 3         | 3         | 0        | CUDA, OMP, OCL          | 312             | –          | 0/3 (0%)           |
| <b>Total</b> | <b>35</b> | <b>96</b> | <b>87</b> | <b>9</b> |                         | <b>80–3,304</b> | <b>271</b> | <b>24/96 (25%)</b> |

All three expose provider-specific reasoning modes, configured at provider-recommended settings (enable\_thinking for Qwen 3.5, reasoning\_effort=medium for GPT-5.4 and GPT-5.3-codex). Because these thinking modes differ in mechanism, observed performance gaps may reflect reasoning-mode effects in addition to base capability. Detailed model configurations, hardware specifications, and cost accounting are provided in Appendix B.4.

**Translation Protocol and Scope.** From the 96-spec corpus (Section 3), we exclude 9 specifications classified as KNOWN\_FAIL due to pre-existing toolchain incompatibilities (Appendix B.4), leaving 87 verified-PASS specs that yield 142 unique translation tasks across ten directions: six standard bidirectional directions among CUDA, OpenMP, and OpenCL, plus four OMP-target case-study directions. Translations follow the kernel-centric protocol (Section 2): only kernel source files are translated while host and build infrastructure remain fixed.

We evaluate oracle-defined translation correctness only (Section 2.4); performance speedup is out of scope.

**Evaluation Phases.** The canonical campaign (L0) evaluates the 142 non-KNOWN\_FAIL pairs from unmodified source code for each model, drawing  $n = 3$  independent samples per task in single-attempt mode (no iterative self-repair). Sampling uses temperature 0.7 for Qwen 3.5; for GPT-5.4 and GPT-5.3-codex, temperature is provider-controlled. Across three models this produces 1,278 L0 records. The augmentation campaign applies an L0-conditional filter, retaining only pairs where at least one L0 sample passes, and evaluates  $n = 1$  sample at each augmentation level L1–L4. Extended sampling configurations and deterministic seed derivation are described in Appendix B.4.

**Metrics.** The primary metric is **pass@ $k$** , using the unbiased estimator from Chen et al. [2021] with normal-approximation confidence intervals on the task-level mean. Separately, we report **per-record pass rates** treating each record as an independent Bernoulli trial, computed with Wilson 95% confidence intervals. Full metric derivations are in Appendix B.4.

**Oracle strength.** PASS denotes successful build, execution, and satisfaction of the declared verification oracle—not a proof of full semantic equivalence. Of the 87 eval-eligible specs, 2 use file-hash verification, 5 use numeric comparison, and the remaining 80 use stdout-pattern plus exit-code checks. We therefore interpret pass rates as *declared-oracle correctness* and report oracle limitations explicitly in Section 6 and Appendix H.

With the methodology established, Section 5 presents the pass rates, failure taxonomy, direction asymmetry, and augmentation robustness results.

## 5 Results

We first report the primary task-level pass@ $k$  results on the balanced L0 task set, then summarize record-level outcomes across the full campaign.

### 5.1 pass@ $k$ Analysis

Restricting to the 142 L0 tasks, Qwen 3.5 achieves pass@1 = 23.9% and pass@3 = 35.2%, while GPT-5.4 achieves pass@1 = 62.7% and pass@3 = 69.7%. GPT-5.3-codex matches GPT-5.4’s pass@1 exactly (62.7%) with a slightly narrower pass@1-to-pass@3 gap (5.6 pp versus 7.0 pp for GPT-5.4; Appendix E.4). Stochastic resampling provides diminishing returns: for Qwen 3.5, 64.8% of tasks

produce no passing sample across three attempts, while only 13.4% pass all three. GPT-5.4 and GPT-5.3-codex show the inverse profile—53.5% and 57.0% pass deterministically—yet 30.3% and 31.7% still hard-fail on every attempt. Across all models, failures are predominantly systematic (incorrect API-surface mapping) rather than stochastic; resampling cannot rescue them.

## 5.2 Record-Level Diagnostic Outcomes

Kernel-centric evaluation isolates executable kernel translation from repository reconstruction. Across all valid records (L0 and L0-conditional augmentation), Qwen 3.5 passes 36.7% (230/626), GPT-5.4 passes 75.5% (621/822), and GPT-5.3-codex passes 74.2% (604/814; full breakdown in Appendix Table 10). GPT-5.3-codex is statistically indistinguishable from GPT-5.4 on the balanced L0 task comparison (Bonferroni-corrected  $p = 1.0$ ,  $OR = 0.93$  [0.74, 1.16], Cohen’s  $h = -0.031$ ); this should not be interpreted as evidence that HPC-specific fine-tuning is unnecessary. Prior work such as ParEval-Repo reports 0% pass@1 on repository-level applications larger than 133 SLoC Davis et al. [2025]; 89% of our kernels exceed that threshold, yet PARBENCH’s kernel isolation enables Qwen 3.5 to achieve 40.3% on CUDA-to-OpenMP. This contrast suggests that repository reconstruction is a major confound in full-application benchmarks, though the different corpora, oracles, and models used in each study preclude a direct comparison.

## 5.3 Direction Dependence

Per-record pass rates vary strongly across directions (Appendix Table 11), reflecting structural API asymmetries. The four OMP-target rows are exploratory case studies ( $m \leq 8$  kernels,  $n \leq 24$  records) and should not be compared directly to the six standard directions; their wide confidence intervals preclude strong conclusions about OMP-target translation in general. Translating from CUDA to OpenMP replaces explicit device memory and launch configuration with compiler directives. The reverse requires introducing explicit synchronization and memory management. OpenCL translations introduce further complexity with separate host/kernel source and JIT compilation, leading to 0% success for OpenCL-to-CUDA under Qwen 3.5. GPT-5.3-codex’s direction profile closely tracks GPT-5.4 across five of six standard directions (maximum gap 7 pp, with two directions matching exactly). The exception is OMP-to-OpenCL, where GPT-5.3-codex outperforms GPT-5.4 by 9.9 pp (82.4% versus 72.5%). Overall, direction difficulty is predominantly a property of the translation task, not the model. At the kernel level, performance is highly heterogeneous (full breakdown in Appendix E.4). Simple kernels pass at high rates, while those with irregular memory access or multi-file dependencies frequently fail entirely.

## 5.4 Failure Taxonomy

Figure 3 places all three models side-by-side on the same scale. Build-stage adaptation is the dominant failure mode: for Qwen 3.5, BUILD\_FAIL accounts for 39.1% of all records and 61.9% of failures; these are incomplete API-surface adaptations (e.g., retained CUDA identifiers in OpenMP targets). GPT-5.4 and GPT-5.3-codex reduce build failures to 15.0% and 17.1% respectively, with residual failures concentrated in OCL→CUDA and OMP→CUDA directions. VERIFY\_FAIL accounts for 4.6% of Qwen 3.5 records, proving that compile-only or zero-exit-code checks are insufficient for rigorous evaluation. RUN\_FAIL (19.3% for Qwen 3.5) is frequently driven by OpenCL JIT compilation errors at runtime.

## 5.5 Augmentation Robustness

To assess whether baseline success is driven primarily by surface-form memorization, we evaluated augmented source variants (L1–L4) on L0-conditional subsets (detailed in Appendix E.4). GPT-5.3-codex shows the same plateau pattern as GPT-5.4: 85.6%–88.7% at L1–L4, in contrast to Qwen 3.5’s monotonic decline from 74.0% at L1 to 56.0% at L4 on the same L0-conditional design. This stability across both GPT models is compatible with robustness to the tested surface-form perturbations. These results are descriptive rather than confirmatory given the L0-conditional design; survivorship bias prevents ruling out deeper structural memorization entirely.

Per-kernel analysis reveals exceptions: on `stencil1d`, Qwen 3.5 produces a passing parallel loop at L0 but reverts to a non-portable CUDA-style tiling pattern at L1 (BUILD\_FAIL) and




Figure 2: Per-kernel translation status across all directions and three models (L0, first sample per task, 35 kernels × up to 8 directions). All-grey rows denote kernels with only KNOWN\_FAIL or phantom specs.




Figure 3: Failure taxonomy by translation direction across all three models (L0, first sample per task). The figure includes only the six standard directions, spanning 120 non-KNOWN\_FAIL tasks in total (24, 24, 19, 19, 17, and 17 tasks by direction), with OMP-target case-study directions omitted for cross-model comparability. Build-stage failures dominate Qwen 3.5 across nearly all directions; GPT-5.4 and GPT-5.3-codex show substantially higher pass counts with residual build failures concentrated in OpenCL→CUDA and OMP→CUDA.

L2 (VERIFY\_FAIL), recovers at L3 (PASS), then introduces a boundary-condition error at L4 (VERIFY\_FAIL)—illustrating how surface-form perturbation can expose fragile pattern-matching even when the model occasionally recovers.

## 6 Discussion, Limitations, and Conclusion

PARBENCH shows that, when repository reconstruction is fixed, current LLMs can produce declared-oracle-passing translations for a non-trivial fraction of kernel-level parallel API tasks, but success remains highly structured by API direction, failure stage, and model tier. The primary bottleneck is

build-stage API adaptation, especially when translations must introduce explicit device-memory and runtime structure. A record-level descriptive omnibus test on 1,278 balanced L0 records ( $\chi^2(2) = 170.43$ ,  $p < 10^{-37}$ ) detects significant heterogeneity; pairwise analysis reveals two tiers: GPT-5.4 and GPT-5.3-codex achieve identical pass rates (267/426; Fisher’s  $p = 1.0$ ), while both surpass Qwen 3.5 by  $5.3\times$  odds (Cohen’s  $h = 0.80$ ). Because the providers use unmatched sampling conditions (Section B.4), the Qwen–GPT gap cannot be causally attributed to model capability alone; the within-provider comparison controls for this confound. Full pairwise statistical details appear in Appendix E.4.

The tier gap is one of determinism, not marginal accuracy: GPT-family models either solve a task reliably across all three samples or fail every attempt, whereas Qwen’s failures are equally systematic but far more frequent (Section 5.1). All three models share the same dominant bottleneck—incomplete API-surface adaptation at build time—but the GPT family resolves the majority of these, leaving residual build failures concentrated in the hardest directions (Section 5.4). Direction rankings are consistent across models: removing explicit parallel constructs (CUDA-to-OpenMP) is substantially easier than introducing them (OpenCL-to-CUDA), suggesting that the difficulty gradient tracks API-surface complexity rather than model-specific memorization (Section 5.3). Augmentation robustness reinforces this interpretation: GPT-family pass rates are stable across L1–L4, while Qwen’s monotonic decline suggests greater reliance on surface-form patterns in the source (Section 5.5).

**Limitations.** The benchmark is correctness-focused: it does not measure speedups, occupancy, memory bandwidth, or efficiency. A harness-passing translation may still be unusably slow, a concern also raised by TRACY Gong et al. [2025]. The corpus is finite; some directions have small paired samples. The L0-conditional augmentation subset ranges from 3–22 kernels per model per standard direction (3–8 for `omp_target` directions), with the smallest cells reflecting Qwen 3.5’s low L0 pass rate, but the balanced cross-model comparison is limited to 12 common kernels in a single direction. Small per-level sample sizes and the L0-conditioned selection of augmentation subsets complicate formal trend interpretation. Ten specs across six kernels were downgraded from strong or medium oracles to weak (stdout pattern plus exit code): six due to cross-API floating-point reduction-order divergence (cf, hotspot, myocyte) and four due to synthesis asymmetry between reference implementations (bfs, nw, nn). This limits oracle strength but correctly avoids penalizing faithful translations. All results are from a single platform (NVIDIA RTX 4070, Ubuntu 24.04, HPC SDK 24.3); for GPU and offload code, compiler and runtime behavior varies across NVIDIA architecture generations and toolchain versions, so results are reproducible on this pinned platform but cross-platform generality remains future work. The stochastic evaluation (temperature 0.7, three samples) captures variability but does not measure iterative repair or agentic translation strategies. The three models use different reasoning mechanisms and sampling configurations (Section B.4). The Qwen–GPT gap is large (Cohen’s  $h = 0.80$ ) but cannot be fully attributed to model capability given unmatched sampling conditions. The within-provider GPT-5.4–GPT-5.3-codex comparison controls for provider-side differences but both share the same Azure OpenAI infrastructure, limiting generalizability beyond this provider. GPT-5.3-codex’s code specialization targets general software engineering tasks; the null result for parallel translation may not generalize to models specifically fine-tuned on HPC code. Appendix H provides a structured Evaluation Card and reporting protocol; the artifact (Appendix I) documents per-record outputs that support reproduction and extension.

## References

- Tomer Bitan, Tal Kadosh, Erel Kaplan, Shira Meiri, Le Chen, Peter Morales, Niranjan Hasabnis, and Gal Oren. UniPar: A unified LLM-based framework for parallel and accelerated code translation in HPC. *arXiv preprint arXiv:2509.12136*, 2025. doi: 10.48550/arXiv.2509.12136. URL <https://arxiv.org/abs/2509.12136>.
- Aman Chaturvedi, Daniel Nichols, Siddharth Singh, and Abhinav Bhatele. HPC-Coder-V2: Studying code LLMs across low-resource parallel languages. In *ISC High Performance 2025 Research Paper Proceedings*, pages 1–14. IEEE, 2025. doi: 10.23919/ISC.2025.11017585. URL <https://doi.org/10.23919/ISC.2025.11017585>.
- Shuai Che, Michael Boyer, Jiayuan Meng, David Tarjan, Jeremy W. Sheaffer, Sang-Ha Lee, and Kevin Skadron. Rodinia: A benchmark suite for heterogeneous computing. In *2009 IEEE International Symposium on Workload Characterization (IISWC)*, pages 44–54. IEEE, 2009. doi: 10.1109/IISWC.2009.5306797. URL <https://doi.org/10.1109/IISWC.2009.5306797>.
- Hao Chen, Qian Liu, and Zhong Ming. Memorize or generalize? an empirical study on code LLMs with semantic-preserving code rewriting, 2025a. URL <https://arxiv.org/abs/2503.02296>.
- Le Chen, Arijit Bhattacharjee, Nesreen K. Ahmed, Niranjan Hasabnis, Gal Oren, Tal Kadosh, and Ali Jannesari. OMPGPT: A generative pre-trained transformer model for OpenMP. In *Proceedings of the 30th International European Conference on Parallel and Distributed Computing (Euro-Par)*, volume 14801 of *Lecture Notes in Computer Science*, pages 121–134. Springer, 2024. doi: 10.1007/978-3-031-69577-3\_9. URL [https://doi.org/10.1007/978-3-031-69577-3\\_9](https://doi.org/10.1007/978-3-031-69577-3_9).
- Le Chen, Nesreen K. Ahmed, Mihai Capotă, Ted Willke, Niranjan Hasabnis, and Ali Jannesari. PCEBench: A multi-dimensional benchmark for evaluating large language models in parallel code generation. In *2025 IEEE International Parallel and Distributed Processing Symposium*, pages 546–557. IEEE, 2025b. doi: 10.1109/IPDPS64566.2025.00055. URL <https://doi.org/10.1109/IPDPS64566.2025.00055>.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgens Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. URL <https://arxiv.org/abs/2107.03374>.
- Joshua H. Davis, Daniel Nichols, Ishan Khillan, and Abhinav Bhatele. On the limits of LLM-based repository-level HPC code translation. In *Proceedings of the 54th International Conference on Parallel Processing (ICPP)*, pages 94–103. ACM, 2025. doi: 10.1145/3754598.3754669. URL <https://doi.org/10.1145/3754598.3754669>.
- Matthew T. Dearing, Yiheng Tao, Xingfu Wu, Zhiling Lan, and Valerie Taylor. LASSI: An LLM-based automated self-correcting pipeline for translating parallel scientific codes. In *2024 IEEE International Conference on Cluster Computing Workshops (CLUSTER Workshops)*, pages 136–143. IEEE, 2024. doi: 10.1109/CLUSTERWorkshops61563.2024.00029. URL <https://doi.org/10.1109/CLUSTERWorkshops61563.2024.00029>.
- Mingzhe Du, Anh Tuan Luu, Bin Ji, and See-Kiong Ng. Mercury: An efficiency benchmark for LLM code synthesis. *arXiv preprint arXiv:2402.07844*, 2024. doi: 10.48550/arXiv.2402.07844. URL <https://arxiv.org/abs/2402.07844>.
- Nichole Etienne, Simon Garcia de Gonzalo, and Dorian Arnold. Openmp-annotated code dataset for large language model fine-tuning on parallel programming tasks. *Frontiers in High Performance Computing*, 4:1771927, 2026. doi: 10.3389/fhpcp.2026.1771927. URL <https://doi.org/10.3389/fhpcp.2026.1771927>.

408 Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach,  
409 Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):  
410 86–92, 2021. doi: 10.1145/3458723. URL <https://doi.org/10.1145/3458723>.

411 William Godoy, Pedro Valero-Lara, Keita Teranishi, and Jeffrey S. Vetter. VibeCodeHPC: An  
412 agent-based iterative prompting auto-tuner for HPC code generation using LLMs, 2025. URL  
413 <https://arxiv.org/abs/2510.00031>.

414 William F. Godoy, Pedro Valero-Lara, Keita Teranishi, Prasanna Balaprakash, and Jeffrey S. Vetter.  
415 Large language model evaluation for high-performance computing software development. *Concur-  
416 rency and Computation: Practice and Experience*, 36(26):e8269, 2024. doi: 10.1002/cpe.8269.  
417 URL <https://doi.org/10.1002/cpe.8269>.

418 Zhihao Gong, Zeyu Sun, Dong Huang, Qingyuan Liang, Jie M. Zhang, and Dan Hao. TRACY: Bench-  
419 marking execution efficiency of LLM-based code translation. arXiv preprint arXiv:2508.11468,  
420 2025. URL <https://arxiv.org/abs/2508.11468>.

421 Re’em Harel, Yuval Pinter, and Gal Oren. Learning to parallelize in a shared-memory environment  
422 with transformers. In *Proceedings of the 28th ACM SIGPLAN Annual Symposium on Principles and  
423 Practice of Parallel Programming*, pages 450–452. ACM, 2023. doi: 10.1145/3572848.3582565.  
424 URL <https://doi.org/10.1145/3572848.3582565>.

425 Re’em Harel, Tal Kadosh, Niranjana Hasabnis, Timothy G. Mattson, Yuval Pinter, and Gal Oren.  
426 PragFormer: Data-driven parallel source code classification with transformers. *International  
427 Journal of Parallel Programming*, 53(1):2, 2025. doi: 10.1007/s10766-024-00778-9. URL  
428 <https://doi.org/10.1007/s10766-024-00778-9>.

429 Ali Reza Ibrahimzada, Kaiyao Ke, Mrigank Pawagi, Muhammad Salman Abid, Rangeet Pan, Saurabh  
430 Sinha, and Reyhaneh Jabbarvand. AlphaTrans: A neuro-symbolic compositional approach for  
431 repository-level code translation and validation. *Proceedings of the ACM on Software Engineering*,  
432 2(FSE):2454–2476, 2025. doi: 10.1145/3729379. URL <https://doi.org/10.1145/3729379>.

433 Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik  
434 Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *Proceedings  
435 of the Twelfth International Conference on Learning Representations (ICLR)*, 2024. doi: 10.48550/  
436 arXiv.2310.06770. URL <https://openreview.net/forum?id=VTF8yNQm66>.

437 Zheming Jin and Jeffrey S. Vetter. A benchmark suite for improving performance portability of the  
438 SYCL programming model. In *2023 IEEE International Symposium on Performance Analysis of  
439 Systems and Software (ISPASS)*, pages 325–327. IEEE, 2023. doi: 10.1109/ISPASS57527.2023.  
440 00041. URL <https://doi.org/10.1109/ISPASS57527.2023.00041>.

441 Zheming Jin, Swaroop Pophale, and Keita Teranishi. Enhancing ChatPORT with CUDA-to-SYCL  
442 kernel translation capability. In *Proceedings of the SC ’25 Workshops of the International Confer-  
443 ence for High Performance Computing, Networking, Storage and Analysis*, pages 524–533. ACM,  
444 2025. doi: 10.1145/3731599.3767398. URL <https://doi.org/10.1145/3731599.3767398>.

445 Tal Kadosh, Niranjana Hasabnis, Timothy Mattson, Yuval Pinter, and Gal Oren. Quantifying OpenMP:  
446 Statistical insights into usage and adoption. In *2023 IEEE High Performance Extreme Computing  
447 Conference (HPEC)*, pages 1–7. IEEE, 2023a. doi: 10.1109/HPEC58863.2023.10363459. URL  
448 <https://doi.org/10.1109/HPEC58863.2023.10363459>.

449 Tal Kadosh, Nadav Schneider, Niranjana Hasabnis, Timothy Mattson, Yuval Pinter, and Gal Oren.  
450 Advising OpenMP parallelization via a graph-based approach with transformers. In *OpenMP: Ad-  
451 vanced Task-Based, Device and Compiler Programming (IWOMP 2023)*, volume 14114 of *Lecture  
452 Notes in Computer Science*, pages 3–17. Springer, 2023b. doi: 10.1007/978-3-031-40744-4\_1.  
453 URL [https://doi.org/10.1007/978-3-031-40744-4\\_1](https://doi.org/10.1007/978-3-031-40744-4_1).

454 Tal Kadosh, Niranjana Hasabnis, Prema Soundararajan, Vy A. Vo, Mihai Capotă, Nesreen K. Ahmed,  
455 Yuval Pinter, and Gal Oren. OMPar: Automatic parallelization with ai-driven source-to-source  
456 compilation. *arXiv preprint arXiv:2409.14771*, 2024a. doi: 10.48550/arXiv.2409.14771. URL  
457 <https://arxiv.org/abs/2409.14771>.

458 Tal Kadosh, Niranjan Hasabnis, Vy A. Vo, Nadav Schneider, Neva Krien, Mihai Capotă, Abdul Wasay,  
459 Guy Tamir, Ted Willke, Nesreen K. Ahmed, Yuval Pinter, Timothy G. Mattson, and Gal Oren.  
460 MonoCoder: Domain-specific code language model for HPC codes and tasks. In *2024 IEEE High*  
461 *Performance Extreme Computing Conference*, pages 1–7. IEEE, 2024b. doi: 10.1109/HPEC62836.  
462 2024.10938441. URL <https://doi.org/10.1109/HPEC62836.2024.10938441>.

463 Erel Kaplan, Tomer Bitan, Lian Ghayeb, Le Chen, Tom Yotam, Niranjan Hasabnis, and Gal Oren.  
464 ParaCodex: A profiling-guided autonomous coding agent for reliable parallel code generation  
465 and translation. *arXiv preprint arXiv:2601.04327*, 2026. doi: 10.48550/arXiv.2601.04327. URL  
466 <https://arxiv.org/abs/2601.04327>.

467 Elias Konstantinidis and Yiannis Cotronis. A quantitative roofline model for GPU kernel performance  
468 estimation using micro-benchmarks and hardware metric profiling. *Journal of Parallel and*  
469 *Distributed Computing*, 107:37–56, 2017. doi: 10.1016/j.jpdc.2017.04.002. URL <https://doi.org/10.1016/j.jpdc.2017.04.002>.

471 Shanchao Liang, Spandan Garg, and Roshanak Zilouchian Moghaddam. The SWE-Bench illusion:  
472 When state-of-the-art LLMs remember instead of reason. *arXiv preprint arXiv:2506.12286*, 2025.  
473 URL <https://arxiv.org/abs/2506.12286>.

474 Quazi Ishtiaque Mahmud, Ali TehraniJamsaz, Hung D. Phan, Le Chen, Mihai Capotă, Theodore L.  
475 Willke, Nesreen K. Ahmed, and Ali Jannesari. AutoParLLM: GNN-guided context generation for  
476 zero-shot code parallelization using LLMs. In *Proceedings of the 2025 Conference of the Nations*  
477 *of the Americas Chapter of the Association for Computational Linguistics: Human Language*  
478 *Technologies (Volume 1: Long Papers)*, pages 11821–11841. Association for Computational  
479 Linguistics, 2025. doi: 10.18653/v1/2025.naacl-long.593. URL [https://aclanthology.org/](https://aclanthology.org/2025.naacl-long.593/)  
480 [2025.naacl-long.593/](https://aclanthology.org/2025.naacl-long.593/).

481 Marko Mišić and Matija Dodović. An assessment of large language models for OpenMP-based  
482 code parallelization: A user perspective. *Journal of Big Data*, 11:161, 2024. doi: 10.1186/  
483 s40537-024-01019-z. URL <https://doi.org/10.1186/s40537-024-01019-z>.

484 Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson,  
485 Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In  
486 *Proceedings of the Conference on Fairness, Accountability, and Transparency, FAT\* ’19*, pages  
487 220–229. ACM, 2019. doi: 10.1145/3287560.3287596. URL [https://doi.org/10.1145/](https://doi.org/10.1145/3287560.3287596)  
488 [3287560.3287596](https://doi.org/10.1145/3287560.3287596).

489 Pei Mu, Yifan Zhu, Dongning Ma, and Xipeng Shen. QiMeng-MuPa: Multi-paradigm parallel code  
490 generation with large language models, 2025. URL <https://arxiv.org/abs/2506.11153>.

491 Daniel Nichols, Joshua H. Davis, Zhaojun Xie, Arjun Rajaram, and Abhinav Bhatele. Can large  
492 language models write parallel code? In *Proceedings of the 33rd International Symposium on*  
493 *High-Performance Parallel and Distributed Computing (HPDC)*, pages 281–294. ACM, 2024a.  
494 doi: 10.1145/3625549.3658689. URL <https://doi.org/10.1145/3625549.3658689>.

495 Daniel Nichols, Aniruddha Marathe, Harshitha Menon, Todd Gamblin, and Abhinav Bhatele. HPC-  
496 Coder: Modeling parallel programs using large language models. In *ISC High Performance 2024*  
497 *Research Paper Proceedings*, pages 1–12. IEEE, 2024b. doi: 10.23919/ISC.2024.10528929. URL  
498 <https://doi.org/10.23919/ISC.2024.10528929>.

499 OpenAI. GPT-5.3-codex system card, 2026. URL [https://openai.com/index/](https://openai.com/index/gpt-5-3-codex-system-card/)  
500 [gpt-5-3-codex-system-card/](https://openai.com/index/gpt-5-3-codex-system-card/). Accessed 2026-05-01.

501 Anne Ouyang, Simon Guo, Simran Arora, Alex L. Zhang, William Hu, Christopher Ré, and Azalia  
502 Mirhoseini. KernelBench: Can LLMs write efficient GPU kernels? In *Proceedings of the 42nd*  
503 *International Conference on Machine Learning*, 2025. doi: 10.48550/arXiv.2502.10517. URL  
504 <https://arxiv.org/abs/2502.10517>.

505 Swaroop Pophale, Zheming Jin, and Keita Teranishi. ChatPORT: Fine-tuned LLM for easy code  
506 PORTing. In *OpenMP: Balancing Productivity and Performance Portability*, volume 16123 of *Lec-*  
507 *ture Notes in Computer Science*, pages 197–211. Springer, 2025. doi: 10.1007/978-3-032-06343-4\_  
508 13. URL [https://doi.org/10.1007/978-3-032-06343-4\\_13](https://doi.org/10.1007/978-3-032-06343-4_13).

509 Baptiste Rozière, Marie-Anne Lachaux, Loïc Chausson, and Guillaume Lample. Unsupervised  
510 translation of programming languages. In *Advances in Neural Information Processing Systems 33*  
511 (*NeurIPS*), 2020. doi: 10.5555/3495724.3497454. URL [https://proceedings.neurips.cc/  
512 paper/2020/hash/ed23fbf18c2cd35f8c7f8de44f85c08d-Abstract.html](https://proceedings.neurips.cc/paper/2020/hash/ed23fbf18c2cd35f8c7f8de44f85c08d-Abstract.html).

513 Nadav Schneider, Tal Kadosh, Niranjan Hasabnis, Timothy Mattson, Yuval Pinter, and Gal Oren. MPI-  
514 RICAL: Data-driven MPI distributed parallelism assistance with transformers. In *Proceedings*  
515 *of the SC '23 Workshops of the International Conference on High Performance Computing,*  
516 *Network, Storage, and Analysis*, pages 2–10. ACM, 2023. doi: 10.1145/3624062.3624063. URL  
517 <https://doi.org/10.1145/3624062.3624063>.

518 Nadav Schneider, Niranjan Hasabnis, Vy A. Vo, Tal Kadosh, Neva Krien, Mihai Capotă, Guy  
519 Tamir, Theodore L. Willke, Nesreen K. Ahmed, Yuval Pinter, Timothy G. Mattson, and Gal Oren.  
520 MPIrigen: MPI code generation through domain-specific language models. In *Proceedings of the*  
521 *2024 Workshop on AI for Systems*, pages 1–6. ACM, 2024. doi: 10.1145/3660605.3660944. URL  
522 <https://doi.org/10.1145/3660605.3660944>.

523 Ali TehraniJamsaz, Arijit Bhattacharjee, Le Chen, Nesreen K. Ahmed, Amir Yazdanbakhsh, and  
524 Ali Jannesari. CodeRosetta: Pushing the boundaries of unsupervised code translation for parallel  
525 programming. In *Advances in Neural Information Processing Systems 37 (NeurIPS)*, 2024. doi:  
526 10.52202/079017-3202. URL <https://openreview.net/forum?id=V6hrg409gg>.

527 John R. Tramm and Andrew R. Siegel. RS Bench – a mini-app for the multipole method of cross  
528 section lookups. Technical report, Argonne National Laboratory, 2014. [https://github.com/  
529 ANL-CESAR/RSBench](https://github.com/ANL-CESAR/RSBench).

530 John R. Tramm, Andrew R. Siegel, Tanzima Islam, and Martin Schulz. XS Bench – the development  
531 and verification of a performance abstraction for Monte Carlo reactor analysis. In *Proceedings of*  
532 *the International Conference on the Physics of Reactors (PHYSOR)*, Kyoto, Japan, 2014. URL  
533 <https://www.mcs.anl.gov/papers/P5064-0114.pdf>.

534 Pedro Valero-Lara, Aaron Young, Thomas III Naughton, Christian Engelmann, Al Geist, Jeffrey S.  
535 Vetter, Keita Teranishi, and William F. Godoy. ChatMPI: LLM-driven MPI code generation for  
536 HPC workloads. In *Proceedings of Supercomputing Asia and International Conference on High*  
537 *Performance Computing in Asia Pacific Region*, pages 19–30. ACM, 2026. doi: 10.1145/3773656.  
538 3773659. URL <https://doi.org/10.1145/3773656.3773659>.

539 Yanli Wang, Yanlin Wang, Suquan Wang, Daya Guo, Jiachi Chen, John Grundy, Xilin Liu, Yuchi  
540 Ma, Mingzhi Mao, Hongyu Zhang, and Zibin Zheng. RepoTransBench: A real-world multilingual  
541 benchmark for repository-level code translation. *IEEE Transactions on Software Engineering*, 52  
542 (2):675–690, 2026. doi: 10.1109/TSE.2025.3645056.

543 Yifan Zhang, Yuqi Zhu, Yibo Liu, Jiawei Liu, Ge Li, and Zhi Jin. CodeMorph: Semantic-preserving  
544 code transformation for data leakage mitigation in code LLM evaluation. In *Proceedings of*  
545 *the 47th International Conference on Software Engineering (ICSE), Industry Track*, 2025. URL  
546 <https://arxiv.org/abs/2506.17627>.

## 547 Appendix Contents

548 This appendix provides supporting material for reproducing and interpreting PARBENCH. Appendix A  
 549 expands the related-work positioning. Appendices B–D document the framework, API-selection  
 550 rationale, and benchmark curation process. Appendix E groups the augmentation protocol and  
 551 extended evaluation results, and Appendix F reproduces the translation prompt. Appendices G–H  
 552 cover reproducibility details and the benchmark’s measurement scope. Appendix I describes artifact  
 553 availability. Appendix J holds the remaining dense reference figures and full per-kernel tables.

|     | Section                                                        | Contents                                                                                                                                                                                                                               |
|-----|----------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|     | <b>A</b> Related Work                                          | Extended comparison with general code benchmarks, repository-level translation benchmarks, OpenMP/MPI corpora, LLM-based parallel-code systems, and robustness or efficiency benchmarks. Includes Tables 2–5.                          |
|     | <b>B</b> Framework and Evaluation Details                      | Specification schema, augmentation levels, prompt anonymization, response extraction, KNOWN_FAIL exclusion policy, sampling configuration, and metric definitions.                                                                     |
|     | <b>C</b> API Selection Rationale                               | Full API co-occurrence analysis, kernel-level API-pair coverage, translation-difficulty taxonomy, and rationale for including CUDA, OpenMP, OpenCL, and OpenMP Target while excluding HIP, SYCL, and OpenACC from the main evaluation. |
|     | <b>D</b> Benchmark Kernel Survey                               | Repository inventory, quality-tier classification, HeCBench selection funnel, Rodinia inventory, XSBench/RS-Bench/mixbench profiles, and exclusion log for benchmarks and kernels.                                                     |
| 554 | <b>E</b> Augmentation Protocol and Extended Evaluation Results | Augmentation transform frequencies, qualitative failure case studies, extended pass-rate analyses, direction-level results, and additional statistical details.                                                                        |
|     | <b>F</b> Translation Prompt Template                           | Full system and user prompt templates used for kernel-centric translation, including output-format requirements and anonymization details.                                                                                             |
|     | <b>G</b> Evaluation Cost and Reproducibility                   | Token usage, task counts, wall-clock evaluation time, and practical budgeting notes for reproducing the three model campaigns.                                                                                                         |
|     | <b>H</b> Evaluation Card                                       | Structured benchmark card defining what PARBENCH measures, what it does not measure, valid and invalid claims, assumptions, user guidance, and recommended reporting protocol.                                                         |
|     | <b>I</b> Artifact Availability                                 | Description of the anonymized supplementary release, including curated specifications, the build-run-verify harness, augmentation and evaluation pipelines, per-record result files, licensing, and planned de-anonymized release.     |
|     | <b>J</b> Supplementary Tables and Figures                      | Remaining dense reference material: direction-level pass@ <i>k</i> figures, cross-suite comparisons, and full per-kernel tables for all three model campaigns.                                                                         |

## Figures and Tables Roadmap

The appendix contains many figures and tables because each group serves a different evidentiary role. The roadmap below explains how to read the supplementary floats: some position PARBENCH against prior work, some justify benchmark design choices, some document corpus construction, some support the empirical findings, and others define reproducibility, cost, artifact availability, and reporting scope.

| Evidence block                                       | Figures and tables                                                            | Why they are included                                                                                                                                                                                                                                                                                                                    |
|------------------------------------------------------|-------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Positioning against prior work                       | Appendix A: Tables 2–5                                                        | These tables explain why PARBENCH is not just another code-generation, parallelization, or repository-level benchmark. They separate prior work by task, granularity, verification style, corpus status, and robustness testing, making clear which evaluation gap PARBENCH is designed to fill.                                         |
| Benchmark mechanics and reproducibility              | Appendix B: Listing 1; Table 8                                                | These items make the benchmark executable and reproducible. The listing shows the declarative JSON specification structure, while the table defines the L0–L4 augmentation levels used throughout the robustness analysis.                                                                                                               |
| API-selection evidence                               | Appendix C: Figures 4–7; Figure 8; Table 6                                    | These floats justify the choice of CUDA, OpenMP, OpenCL, and OpenMP Target. They show both repository-level and kernel-level API availability, expose why repository counts alone are misleading, and summarize the programming-model differences that make the selected translation directions technically meaningful.                  |
| Corpus construction and filtering                    | Appendix D: Figures 9–14; Table 7                                             | These floats document how the final benchmark corpus was selected from a much larger HPC benchmark landscape. They show the repository inventory, API breadth, candidate ranking, HeCBench funnel, verification methods, and final corpus characteristics, supporting the claim that the corpus is systematic rather than cherry-picked. |
| Augmentation and memorization probes                 | Appendix E: Figure 15; Table 15; Table 16; Figures 16–17                      | These floats explain the source-perturbation experiment. They show which AST transforms were applied, how often they appear, and whether model success persists under L1–L4 perturbations. Their role is to support the robustness and surface-form-memorization analysis.                                                               |
| Prompt and anonymization protocol                    | Appendix F: Listings 4 and 5                                                  | These listings document the actual translation prompt templates (system and user messages). They are included so that future users can distinguish model capability from prompt-format differences.                                                                                                                                      |
| Overall evaluation setup                             | Appendix G: Tables 18 and 19                                                  | These tables record the model, provider, reasoning-mode, hardware, compiler, and software configuration details needed to interpret the empirical results. They also make explicit where sampling conditions differ across providers.                                                                                                    |
| Pass-rate, direction analysis, and sampling behavior | Appendix E: Table 10; Table 11; Table 12; Table 13; Appendix J: Figures 18–20 | These floats support the record-level pass-rate summary, direction-level analysis, and pass@ $k$ analysis. They distinguish single-sample success from repeated-sampling success and show whether failures are stochastic or systematic across directions and models.                                                                    |
| Kernel-level difficulty                              | Appendix E: Table 14; Main body: Figure 2; Appendix J: Tables 24, 25, and 26  | These floats show that translation difficulty is highly uneven across kernels. They are included to prevent the aggregate pass rates from hiding hard kernels, easy kernels, multi-file effects, and direction-specific failure patterns.                                                                                                |
| Failure taxonomy and direction asymmetry             | Appendix E: Table 9; Table 17                                                 | These floats diagnose where translation breaks. They separate build, runtime, verification, and extraction failures, and they support the paper’s claims about API-surface adaptation, direction asymmetry, and multi-file or host-scaffolding difficulty.                                                                               |
| Cross-suite generality                               | Appendix J: Figures 21–23                                                     | These figures show whether results are concentrated in one benchmark suite or persist across Rodinia, HeCBench, XSBench, RSBench, and mixbench. Their role is to support claims about suite-level generality and domain variation.                                                                                                       |
| Repository-level versus kernel-level opportunity     | Appendix C: Figure 8                                                          | This figure explains one of the benchmark-design motivations: repository-level API co-occurrence substantially underestimates the number of independent kernel-level translation opportunities. It supports the kernel-centric design choice.                                                                                            |
| Evaluation cost and practical reproducibility        | Appendix G: Table 20                                                          | This table reports the practical cost of running the evaluation campaign. It helps future users estimate token usage, task volume, wall-clock time, and budget before reproducing or extending PARBENCH.                                                                                                                                 |
| Benchmark scope and reporting protocol               | Appendix H: Tables 21, 22, and 23                                             | These tables define what PARBENCH measures, what it does not measure, which claims are valid, and what future papers should report. They are included to make the benchmark harder to misuse or overclaim.                                                                                                                               |

## A Related Work

The main text positions PARBENCH through the evaluation gap it addresses. This appendix expands that comparison in Tables 2–5, which organize prior work by evaluation level and role: general code and repository-level benchmarks, OpenMP/MPI and local parallel-code corpora, LLM-based parallel-code translation and porting systems, and efficiency or robustness benchmarks. The comparison separates five questions that are often conflated in prior work: whether the task is generation, parallelization, porting, or translation; what unit of code is evaluated; whether the corpus is a reusable benchmark or a paper-specific local evaluation set; whether correctness is checked by execution; and whether robustness to surface-form variation is measured. This distinction is important because recent LLM-for-parallel-programming work has produced many useful local corpora, benchmark subsets,

and task-specific evaluation protocols, but comparatively few reusable executable frameworks for controlled cross-API translation of existing parallel kernels.

**Local parallel-code corpora and paper-specific benchmarks.** A recurring pattern in recent LLM-assisted parallel-programming work is that evaluation is often built around a corpus or benchmark subset introduced for the specific paper. Earlier OpenMP and MPI work constructed curated datasets for source-to-source parallelization, pragma recommendation, parallel-region classification, MPI API assistance, and domain-specific model training Harel et al. [2023], Schneider et al. [2023], Kadosh et al. [2023b], Chen et al. [2024], Schneider et al. [2024], Kadosh et al. [2024b], Harel et al. [2025], Etienne et al. [2026]. Prior corpus studies also show that OpenMP usage itself has been systematically characterized at scale Kadosh et al. [2023a], reinforcing that parallel-programming corpora are valuable but do not by themselves define a controlled LLM translation benchmark. Later LLM-based systems for automatic parallelization, porting, and HPC code modeling continued this pattern, evaluating on paper-specific mixtures of OpenMP, MPI, CUDA, mini-applications, or benchmark-suite kernels Kadosh et al. [2024a], Mahmud et al. [2025], Mišić and Dodović [2024], Nichols et al. [2024b], Chaturvedi et al. [2025], Godoy et al. [2024], Valero-Lara et al. [2026], Godoy et al. [2025]. These studies are directly relevant because they show sustained demand for parallel-code evaluation, but they do not provide a shared measurement substrate for controlled cross-API translation of existing parallel kernels.

**Generation, parallelization, porting, and translation measure different capabilities.** Parallel code generation benchmarks such as ParEval and PCEBench ask whether a model can synthesize a parallel program from a natural-language description, and ParEval also includes limited translation settings Nichols et al. [2024a], Chen et al. [2025b]. These are important capabilities, but they differ from source-to-source API translation under fixed executable infrastructure. In generation, the model may choose its own decomposition, data layout, and implementation strategy; in translation, it must preserve the structure and behavior of an existing implementation while adapting thread indexing, synchronization, memory movement, host-device coordination, and API-specific launch structure. Automatic-parallelization systems such as OMPPar, AutoParLLM, and related OpenMP-focused assessments measure yet another capability: identifying and introducing parallelism into serial or partially parallel code Kadosh et al. [2024a], Mahmud et al. [2025], Mišić and Dodović [2024]. Porting frameworks such as ChatPORT and CUDA-to-SYCL translation work are closer to PARBENCH in spirit, but are typically narrower in API direction, corpus scope, or evaluation protocol Pophale et al. [2025], Jin et al. [2025]. PARBENCH instead targets cross-API translation of already-parallel kernels, where the reported claim is satisfaction of the declared oracle for an existing implementation rather than the existence of some correct parallel solution.

**Repository-level realism versus kernel-level attribution.** Repository-level benchmarks expose realistic integration barriers. SWE-bench and RepoTransBench demonstrate the importance of repository-level executable evaluation in general software settings Jimenez et al. [2024], Wang et al. [2026], while ParEval-Repo shows that full-repository HPC translation can be dominated by build-system and scaffolding failures Davis et al. [2025]. Those barriers matter for deployment, but they can obscure the narrower scientific question of whether the model can translate the computational kernel itself. PARBENCH occupies the intermediate granularity between single-function benchmarks and full repositories. It uses real kernels from established HPC suites, but fixes the surrounding build, run, and verification infrastructure through declarative specifications. This design preserves executable evaluation while making failure attribution sharper: a failed task is more directly tied to API adaptation, multi-file kernel coordination, or semantic preservation, rather than to missing scaffolding or repository reconstruction.

**Agentic repair and specialized models are complementary to benchmark design.** Several systems improve parallel-code generation or translation through fine-tuning, prompting, or iterative agentic repair. CodeRosetta and HPC-Coder-style models demonstrate the value of domain-specific corpora and model adaptation for HPC and parallel code TehraniJamsaz et al. [2024], Nichols et al. [2024b], Chaturvedi et al. [2025]. LASSI evaluates an automated self-correcting pipeline for parallel scientific-code translation on a curated HeCBench subset Dearing et al. [2024]. UniPar, QiMeng-MuPa, VibeCodeHPC, and ParaCodex similarly explore multi-paradigm generation, sequential-to-parallel transformation, agentic HPC code generation, or profiling-guided repair using local or mixed benchmark suites Bitan et al. [2025], Mu et al. [2025], Godoy et al. [2025], Kaplan et al. [2026].

Table 2: Representative related work, Part I-a. General code benchmarks and repository-level translation benchmarks.

| Work                                 | Primary task / granularity                           | Corpus and verification basis                            | Relation to PARBENCH                                                                        |
|--------------------------------------|------------------------------------------------------|----------------------------------------------------------|---------------------------------------------------------------------------------------------|
| HumanEval Chen et al. [2021]         | General code generation; function level              | Public benchmark with unit tests                         | Establishes execution-based evaluation, but does not exercise parallel API semantics.       |
| TransCoder Rozière et al. [2020]     | General source-to-source translation; function level | Public multilingual corpus with unit tests               | Provides a code-translation precedent, but not for HPC or parallel APIs.                    |
| SWE-bench Jimenez et al. [2024]      | Repository issue repair; repository level            | Public GitHub-issue benchmark with repository tests      | Demonstrates executable repository evaluation, but is mostly non-HPC and non-parallel.      |
| RepoTransBench Wang et al. [2026]    | Repository-level code translation                    | Public repository benchmark with build/test signals      | Shows repository translation complexity, but confounds kernel translation with scaffolding. |
| AlphaTrans Ibrahimzada et al. [2025] | Repository translation and validation                | Repository-scale translation tasks with validation/tests | Adjacent repository-level translation work, but not focused on parallel API semantics.      |

Table 3: Representative related work, Part I-b. OpenMP/MPI/local parallel-code corpora and generation benchmarks.

| Work                                           | Primary task / granularity                                             | Corpus and verification basis                                    | Relation to PARBENCH                                                                                                                          |
|------------------------------------------------|------------------------------------------------------------------------|------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| Learning to Parallelize Harel et al. [2023]    | OpenMP parallelization prediction; loop/region level                   | Curated OpenMP training and evaluation corpus                    | Shows local corpus construction for OpenMP assistance, not executable cross-API translation.                                                  |
| MPI-RICAL Schneider et al. [2023]              | MPI assistance; function/region level                                  | Curated MPI-focused corpus with recommendation metrics           | Relevant as MPI-specific parallel-code assistance using a local dataset.                                                                      |
| OpenMP Graph Transformer Kadosh et al. [2023b] | OpenMP parallelization advice; loop/region level                       | Curated OpenMP corpus with recommendation/classification metrics | Relevant to pragma advice, but not translation evaluation.                                                                                    |
| OMPGPT Chen et al. [2024]                      | OpenMP-oriented code modeling; function/task level                     | OpenMP-oriented model corpus and generation metrics              | Domain-specific modeling work, but not reusable translation infrastructure.                                                                   |
| MPIrigen Schneider et al. [2024]               | MPI code generation; function/task level                               | MPI-specific curated corpus with generation checks               | Generation-oriented MPI work, not API-to-API kernel translation.                                                                              |
| MonoCoder Kadosh et al. [2024b]                | HPC code modeling; task/function level                                 | HPC-oriented model corpus and task metrics                       | Relevant HPC specialization work, but method/corpus-specific.                                                                                 |
| PragFormer Harel et al. [2025]                 | Parallel-code classification; loop/region level                        | Curated classification corpus                                    | Relevant representation work, but not executable translation.                                                                                 |
| OMPPar Kadosh et al. [2024a]                   | AI-driven OpenMP source-to-source parallelization; program/loop level  | Paper-specific OpenMP evaluation set                             | Automatic parallelization rather than translation of existing parallel kernels.                                                               |
| AutoParLLM Mahmud et al. [2025]                | Zero-shot code parallelization; loop/function level                    | Local parallelization evaluation set                             | Measures introducing parallelism, not preserving existing parallel semantics across APIs.                                                     |
| OpenMP Assessment Mišić and Dodović [2024]     | LLMs for OpenMP parallelization; snippet/program level                 | Paper-specific assessment set                                    | Shows the need for OpenMP-specific evaluation, but not cross-API translation.                                                                 |
| ParEval Nichols et al. [2024a]                 | Parallel code generation with limited translation settings; task level | Public benchmark with correctness checks                         | Includes generation and some translation settings, but not PARBENCH’s fixed executable kernel corpus or broader cross-API direction coverage. |
| PCEBench Chen et al. [2025b]                   | Multi-dimensional parallel code generation; task level                 | Public benchmark with correctness/task metrics                   | Complements PARBENCH by evaluating generation rather than existing-kernel translation.                                                        |

Table 4: Representative related work, Part II. LLM-based parallel-code translation, porting, and HPC-code modeling systems.

| Work                                                                                  | Primary task / granularity                                                                                       | Corpus and verification basis                                                                                  | Relation to PARBENCH                                                                                                                                  |
|---------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|
| LASSI Dearing et al. [2024]                                                           | Agentic parallel-code translation; kernel level                                                                  | Curated HeCBench subset with build-run-verify                                                                  | Closest in verification style, but smaller in scope and focused on iterative repair.                                                                  |
| CodeRosetta TehraniJamsaz et al. [2024]                                               | Parallel code translation and fine-tuning; function/kernel level                                                 | Model-specific C++/CUDA corpus with compile/similarity/local checks                                            | Relevant translation work, but lacks fixed multi-API executable specifications.                                                                       |
| HPC-Coder Nichols et al. [2024b]<br>HPC-Coder-V2 Chaturvedi et al. [2025]             | HPC code modeling; task/function level<br>Code LLMs across low-resource parallel languages; task level           | HPC-oriented corpus and model-evaluation metrics<br>Parallel-language tasks with pass@k and task checks        | Domain-specific HPC modeling, not benchmark-infrastructure focused. Relevant to parallel-code LLMs, but not controlled executable kernel translation. |
| ParEval-Repo Davis et al. [2025]                                                      | Repository-level HPC translation                                                                                 | Public HPC repository benchmark with build and functional tests                                                | Motivates PARBENCH by showing that repository integration can dominate failures.                                                                      |
| UniPar Bitan et al. [2025]                                                            | Parallel and accelerated code translation; task/program level                                                    | System-specific evaluation corpus with local correctness checks                                                | Relevant multi-paradigm translation, but not primarily a reusable benchmark substrate.                                                                |
| QiMeng-MuPa Mu et al. [2025]                                                          | Sequential-to-parallel translation; task/program level                                                           | Curated training/evaluation corpus with local correctness checks                                               | Related parallelization/translation work, but different from API-to-API translation of existing kernels.                                              |
| ChatPORT Pophale et al. [2025]<br>ChatPORT<br>CUDA-to-SYCL Jin et al. [2025]          | LLM-assisted code porting; kernel/program level<br>CUDA-to-SYCL translation; kernel level                        | Local porting evaluation with build/run checks<br>Local CUDA-to-SYCL evaluation with build/run checks          | Relevant porting work, but method-specific.<br>Directly relevant API migration, but narrower in direction and scope.                                  |
| ChatMPI Valero-Lara et al. [2026]<br>Large LLM Evaluation for HPC Godoy et al. [2024] | MPI code generation; function/workload level<br>Evaluating LLMs for HPC software development; task/program level | MPI-focused local evaluation with functional checks<br>HPC-oriented evaluation suite with task-specific checks | Distributed-parallel generation, not cross-API kernel translation.<br>Broad HPC LLM context, but not controlled cross-API translation.                |
| VibeCodeHPC Godoy et al. [2025]                                                       | Agentic HPC code generation/tuning; application/kernel level                                                     | Small local HPC benchmark set with build/run/performance signals                                               | Agentic generation and optimization, not existing-kernel API translation.                                                                             |
| ParaCodex Kaplan et al. [2026]                                                        | Profiling-guided agentic parallel generation/translation; kernel/program level                                   | Curated suite mixture with build-run-verify and profiling feedback                                             | Directly relevant method-side work; PARBENCH can evaluate such systems under fixed specs.                                                             |

Table 5: Representative related work, Part III. Efficiency, robustness, and benchmark-validity studies relevant to PARBENCH.

| Work                                                         | Primary task / granularity                                                                                                   | Corpus and verification basis                                                                         | Relation to PARBENCH                                                                                                                                       |
|--------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| TRACY Gong et al. [2025]                                     | Efficiency of LLM-based code translation; function/task level                                                                | Efficiency benchmark with correctness and runtime                                                     | Complementary performance axis; PARBENCH focuses first on correctness.                                                                                     |
| KernelBench Ouyang et al. [2025]<br>Mercury Du et al. [2024] | Efficient GPU-kernel generation; kernel level<br>Efficiency of LLM code synthesis; function/task level                       | GPU-kernel benchmark with correctness and speed<br>Efficiency benchmark with runtime/resource metrics | Accelerator-kernel generation, not translation of existing parallel APIs. Evaluates efficiency rather than parallel API semantic preservation.             |
| CodeMorph Zhang et al. [2025]                                | Robustness and leakage via code transformation; function level                                                               | Existing benchmarks with transformed variants and unit tests                                          | Provides the source-perturbation principle adapted by PARBENCH to HPC translation.                                                                         |
| Semantic Code Rewrite Chen et al. [2025a]                    | Memorization versus generalization under code rewriting; function level                                                      | Rewritten code benchmarks with unit tests                                                             | Motivates source rewriting as a contamination and robustness probe.                                                                                        |
| SWE-bench Illusion Liang et al. [2025]<br><b>PARBENCH</b>    | Benchmark memorization analysis; repository level<br><b>Cross-API translation of existing parallel kernels; kernel level</b> | SWE-bench-style evaluation<br><b>Reusable executable specs with conjunctive build-run-verify</b>      | Supports robustness checks for public code benchmarks.<br><b>Fixes infrastructure, isolates parallel API translation, and adds L0-L4 AST augmentation.</b> |

These systems answer questions about method design: how far can feedback, fine-tuning, or agentic iteration push performance? PARBENCH is complementary: it provides an evaluation substrate on which such methods can be compared under a controlled kernel-centric translation protocol, including first-attempt evaluation, fixed infrastructure, and uniform failure taxonomy.

**Verification rigor.** Verification practices vary substantially across the literature. Some works report prediction accuracy, pragma classification, compilation success, similarity, or  $\text{pass}@k$  on task-specific checks Harel et al. [2023], Kadosh et al. [2023b], TehraniJamsaz et al. [2024], Chaturvedi et al. [2025]. Others incorporate build, run, and functional validation, especially in agentic or repository-level settings Dearing et al. [2024], Davis et al. [2025], Kaplan et al. [2026]. PARBENCH adopts conjunctive build-run-verify evaluation at the kernel level: a translation must compile, execute, and satisfy all declared oracle checks. This matters because parallel translations can compile and even terminate successfully while still violating output constraints, omitting required host-device transfers, mishandling reductions, or preserving only superficial API structure. Compile-only or exit-code-only evaluation would therefore overestimate translation capability.

**Efficiency, performance reasoning, and executable validity.** A separate line of work evaluates whether generated or translated code is efficient rather than merely passing task oracles. TRACY, KernelBench, Mercury, and related GPU-kernel or efficiency benchmarks study runtime, kernel quality, optimization, or resource behavior Gong et al. [2025], Ouyang et al. [2025], Du et al. [2024]. These works are complementary because declared-oracle PASS and performance are separable objectives: a translation can pass the declared oracle but be slow, or run fast while still violating the intended transform. PARBENCH deliberately scopes to executable validity under declared oracles in order to isolate the translation-capability question before adding performance tuning as a second axis.

**Robustness and benchmark validity.** Public code benchmarks face contamination and memorization risks, and recent work shows that surface-form perturbations can reveal brittle pattern matching Liang et al. [2025], Zhang et al. [2025], Chen et al. [2025a]. CodeMorph applies intended semantics-preserving transformations to standard code benchmarks to diagnose leakage and robustness. PARBENCH applies the same principle to HPC translation, where contamination risk is especially plausible because suites such as Rodinia and HeCBench appear in public repositories and prior papers. Its AST-driven augmentation engine creates L0–L4 source variants using behavior-preserving transformations such as identifier renaming, arithmetic rewriting, pointer/array interchange, typedef expansion, and condition rewriting, with baseline execution used to validate the retained subset. The goal is not only to detect possible memorization, but to measure whether translation success survives controlled perturbations of the source surface form.

Appendix H provides a structured Evaluation Card declaring PARBENCH’s measurement scope, assumptions, valid and invalid claims, and a recommended reporting protocol for future users.

PARBENCH is evaluation infrastructure rather than a new translation method. Its contribution is to make the measurement problem sharper: kernel-centric isolation avoids repository-level build confounds, declarative specs make declared-oracle evaluation reproducible, augmentation probes robustness, and the harness exposes where translations fail. The initial results show that current LLMs have real but uneven parallel translation capability, with build adaptation, direction asymmetry, and multi-file coordination as central barriers to reliable use. The benchmark code, all 96 curated JSON specs (87 eval-eligible), and per-record result files for all three models are included in the supplementary material to enable independent verification and extension.

## B Framework and Evaluation Details

This appendix expands the framework description in Section 2. It records the spec structure, augmentation level definitions, and implementation details that are useful for reproducing the benchmark but too detailed for the main narrative.

```

{
  "identity": {
    "unique_id": "rodinia-bfs-cuda",
    "parallel_api": "cuda",
    "source_suite": "rodinia"
  },
  "files": {
    "prompt_payload": ["bfs.cu", "kernel.cu", "kernel2.cu"],
    "support_files": ["Makefile"],
    "translation_targets": ["bfs.cu", "kernel.cu", "kernel2.cu"]
  },
  "build": {
    "commands": {"build": "make CUDA_DIR=..."},
    "outputs": {"executable": "./bfs.out"}
  },
  "run": {
    "executable": "./bfs.out",
    "timeout_seconds": 300,
    "input_configurations": {
      "correctness": {
        "arguments": ["../data/bfs/graph1MW_6.txt"]
      }
    }
  },
  "verification": {
    "strategies": [
      {"type": "stdout_pattern",
       "pattern": "Result stored in result\\.\\.txt"},
      {"type": "exit_code", "expected": 0}
    ]
  }
}

```

Listing 1: Condensed specification structure for rodinia-bfs-cuda.

## 674 B.1 Specification Schema

675 Each benchmark kernel-API is encoded by a JSON specification. The schema defines five primary  
676 field groups (with additional implementation, hardware, metadata, baseline\_results, and  
677 performance fields defined in the schema):

- 678 • **Identity and provenance:** globally unique identifier, source suite, kernel name, target API,  
679 repository URL, and commit.
- 680 • **File partitioning:** prompt payload files, translation targets, support files, and verification-  
681 only files.
- 682 • **Build:** build system, clean/configure/build commands, environment variables, and expected  
683 executable path.
- 684 • **Run:** executable, input configuration, command-line arguments, timeouts, and optional  
685 environment variables.
- 686 • **Verification:** ordered verification strategies applied conjunctively. Five types  
687 are implemented: `exit_code`, `stdout_pattern`, `stdout_exclude_pattern`,  
688 `numeric_comparison`, and `file_hash`.

## 689 B.2 Prompt Anonymization

690 Before prompt construction, PARBENCH applies six anonymization measures: it removes the kernel  
691 name and description, strips C/C++ comments from all files shown to the LLM, replaces source file-  
692 names with generic labels (Source File 1, etc.), genericizes target filenames (`translated_0.ext`,  
693 etc.) and support-file names (Header File *N*, Code File *N*, Infrastructure File *N*), and  
694 anonymizes kernel identifiers in build commands. These measures complement source augmentation:  
695 anonymization reduces benchmark identity leakage, while augmentation changes the source surface  
696 form. The complete prompt template is reproduced in Appendix F.

### 697 B.3 Response Extraction Strategy

698 Each LLM response is parsed to recover the translated source files using a multi-tier extraction  
699 strategy that progressively relaxes matching criteria. The tiers are applied in order until all expected  
700 target files are recovered:

- 701 1. **Explicit filename match.** The parser searches for markdown code fences annotated with  
702 the expected target filename (e.g., `““cpp filename=translated_0.cpp”`). This is the  
703 highest-confidence tier.
- 704 2. **Extension-based match.** If explicit filenames are absent, the parser matches code fences by  
705 file extension (e.g., `.cu`, `.cl`, `.cpp`) against the expected target file types.
- 706 3. **Fuzzy match.** Partial filename matches and common naming variants are attempted (e.g.,  
707 matching `kernel.cu` to an expected `translated_0.cu`).
- 708 4. **Elimination-based assignment.** When multiple code blocks remain unmatched, they are  
709 assigned to the remaining expected files by elimination order.

710 If any expected target file cannot be recovered or the extracted content is empty after all tiers, the  
711 task is classified as `EXTRACT_FAIL`. Across the 2,262 evaluation-eligible records, extraction failures  
712 are rare (1 for Qwen 3.5, 3 for GPT-5.4, 0 for GPT-5.3-codex), indicating that the structured prompt  
713 format (Appendix F) effectively guides models to produce parseable output.

### 714 B.4 Extended Experimental Setup Details

715 This section expands on the experimental setup described in Section 4.

716 **KNOWN\_FAIL Exclusion Policy.** Nine specs are classified as `KNOWN_FAIL` due to toolchain  
717 incompatibilities or pre-existing runtime failures unrelated to LLM translation quality: seven Ro-  
718 dinia specs affected by deprecated CUDA 12 texture APIs, missing system libraries, pre-existing  
719 segmentation faults, or silent OpenCL runtime failures, and two HeCBench `omp_target` specs with  
720 build or verification failures on the evaluation platform. A record is excluded when *either* the source  
721 or the target spec is `KNOWN_FAIL`. Source-side exclusion is necessary because source baseline  
722 validity is unverified on the failing platform; target-side exclusion is necessary because the target  
723 build or verification infrastructure is broken. For Qwen 3.5, this policy excludes 82 of 708 total  
724 records, leaving 626 evaluation-eligible records (426 L0 + 200 augmentation records). The GPT-5.4  
725 and GPT-5.3-codex evaluation batches pre-exclude `KNOWN_FAIL` specs, so all 822 GPT-5.4 and  
726 814 GPT-5.3-codex on-disk records are evaluation-eligible.

727 **Sampling Configuration.** For Qwen 3.5, we set temperature 0.7 with  $\text{top-}p = 1.0$  (full-distribution  
728 sampling). For GPT-5.4 and GPT-5.3-codex, Azure API does not permit explicit temperature or  
729  $\text{top-}p$  parameters on reasoning-model deployments—these are controlled internally by the provider  
730 and cannot be overridden by the caller. Cross-model differences in pass rates may therefore partly  
731 reflect unmatched sampling conditions rather than model capability alone; this confound is imposed  
732 by provider API constraints, not benchmark design, and we report it transparently. A deterministic  
733 seed derived from SHA-256 of the concatenation `source_spec|target_spec|sample_id` (pipe-  
734 delimited, truncated to 31 bits) is passed to Qwen 3.5 (Together AI) and GPT-5.4 (Azure Chat  
735 Completions). The Responses API used by GPT-5.3-codex does not accept a seed parameter. Seeds  
736 are recorded for auditing but do not guarantee bitwise-identical outputs: Together AI treats seeds as  
737 best-effort for MoE routing, and Azure provides no determinism guarantee even with a fixed seed.  
738 Each sample receives a single LLM call with no iterative feedback, isolating first-attempt translation  
739 capability.

740 **Canonical Evaluation and Augmentation.** The canonical campaign evaluates all 142 unique  
741 pairs at augmentation level L0 (unmodified source code) with three samples each, producing 426 L0  
742 records per model and 1,278 total. Each record passes through the full build–run–verify pipeline with  
743 conjunction semantics. Most specs verify via exit-code and stdout-pattern checks; five additionally  
744 declare numeric-comparison oracles for floating-point outputs, and two verify via exit-code and  
745 file-hash oracles instead. For the augmentation, an L0-conditional filter is applied: a pair is included  
746 only if any of its three canonical L0 samples passes. This avoids spending budget on pairs the model

cannot translate even from unmodified source. For Qwen 3.5, 50 of the 142 pairs qualify (35%); for GPT-5.4, 99 pairs qualify (70%); for GPT-5.3-codex, 97 pairs qualify (68%). Each qualifying pair is evaluated once ( $k = 1$ ) at each of levels L1 through L4.

**Metrics Details.** For  $\text{pass}@k$ , following Chen et al. [2021], we use the unbiased estimator  $1 - \binom{n-c}{k} / \binom{n}{k}$ , where  $c$  is the number of passing samples out of  $n$  total. We report  $\text{pass}@1$  (expected single-sample success rate, equivalent to  $c/n$  averaged across tasks) and  $\text{pass}@3$  (fraction of tasks where any sample passes). Separately, we report per-record pass rates with Wilson 95% confidence intervals. Wilson intervals are preferred over the Wald interval for their superior coverage at extreme proportions. Since three samples per task share the same kernel and direction, the independence assumption is anti-conservative and the Wilson CI understates the true uncertainty. For augmentation analysis, we apply the Cochran–Armitage trend test where per-level sample sizes are adequate, with Cohen’s  $h$  for adjacent-level effect sizes.

## C API Selection Rationale

This appendix provides the complete quantitative evidence underlying PARBENCH’s selection of CUDA, OpenMP, and OpenCL as target parallel programming APIs. Section 3.1 in the main paper summarizes the corpus selection process; here we present the expanded survey data and kernel-level analysis that informed the final selection.

### C.1 Full API Co-Occurrence Data

The benchmark landscape survey initially covered 71 repositories spanning benchmark suites, mini-applications, proxy applications, libraries, full applications, and microbenchmarks across 29 distinct parallel programming APIs. Of these, 35 met the inclusion criteria for detailed analysis (Section D.2). The main paper (Section 3.1) reports statistics for these 35 qualifying repositories.

Figure 4 shows the full survey. The dominant pattern is a dense co-occurrence cluster among MPI, OpenMP, CUDA, HIP, and OpenACC, with sparser connectivity among portability-layer (Kokkos, RAJA, SYCL) and directive-based (OpenMP Target) APIs. Peripheral APIs such as Chapel, STM/TM, HLS, and Python show near-zero co-occurrence.

We classify the co-occurrence structure into three groupings:

**Core cluster.** MPI, OpenMP, CUDA, HIP, and OpenACC form a densely interconnected cluster in the co-occurrence matrix, with the highest pairwise counts concentrated in suites with broad API coverage (HeCBench, BabelStream, RAJAPerf, CloverLeaf). OpenCL co-occurs with fewer core APIs (present in 13 of the 71 surveyed repositories, versus 38–50 for the top three) but is included in PARBENCH’s target API set because it exercises a qualitatively distinct programming model (Section C.3).

**Portability periphery.** SYCL, Kokkos, RAJA, OpenMP Target, and MPI+OpenMP connect to the core cluster with moderate co-occurrence, concentrated in a smaller number of suites (primarily RAJAPerf, BabelStream, and HeCBench). Pthreads, TBB, and C++ PSTL appear in 3–4 repositories each and occupy a similar peripheral role.

**Isolated APIs.** OCCA, Thrust, Chapel, C++ AMP, HMPP, Python, HLS, GSParLib, STM/TM, Serial, and UPC++ have maximum pairwise co-occurrence at most 2. Coarray Fortran, Charm++, UPC, and OpenSHMEM reach co-occurrence of 3–4 with MPI or with each other but remain disconnected from the GPU-centric core.

Figure 5 presents the per-API coverage in a ranked bar chart, making the quantitative dominance of the top-3 APIs clear.

APIs below 5 benchmarks were excluded from evaluation consideration due to insufficient benchmark material for statistical evaluation.

### Kernel-Level Co-Occurrence

The repository-level analysis in Figures 4–5 counts how many *repositories* provide implementations in a given API pair, but does not capture the depth of coverage within each repository. A suite




Figure 4: API co-occurrence heatmap across all 71 surveyed repositories. Cell values indicate the number of repositories providing implementations in both the row and column APIs.

like HeCBench provides 522 individual kernels, each potentially implemented in multiple APIs, while a single-kernel mini-application like miniBUDE provides only 1 kernel across 12+ APIs. For translation evaluation, the kernel count is the relevant unit: it determines how many independent evaluation instances are available per direction.

Figure 6 provides the strongest quantitative evidence for API selection. The kernel-level data reveals a clear hierarchy:

- **Tier 1 (>600 kernel pairs):** CUDA–HIP (633), CUDA–SYCL (616), HIP–SYCL (615). These pairs are dominated by HeCBench’s systematic multi-API coverage.
- **Tier 2 (~450 kernel pairs):** CUDA–OpenMP (472), HIP–OpenMP (453), SYCL–OpenMP (453). OpenMP’s CPU threading model provides the deepest pool for *paradigm-crossing* translation evaluation.
- **Tier 3 (~100 kernel pairs):** OpenMP Target and Sequential at 106 each, concentrated in RAJAPerf.
- **Tier 4 (<30 kernel pairs):** OpenCL at 27 (CUDA–OpenCL) and 24 (OpenMP–OpenCL), with an OpenCL diagonal of 27 total kernels.




Figure 5: API coverage across the benchmark survey. OpenMP (50 benchmarks), MPI (45), and CUDA (38) dominate, with OpenACC (17), HIP (15), and MPI+OpenMP (15) forming a second tier. OpenCL (13) and Kokkos (12) bridge the second and third tiers. OpenMP Target (8), SYCL (7), and RAJA (7) form a third tier.

810 The network visualization in Figure 7 makes this tiered structure visually apparent.

811 Figure 8 complements the full survey views above by showing why repository-level counts alone  
812 would understate the number of independent translation opportunities available for evaluation.

## 813 C.2 Translation Difficulty Taxonomy

814 The preceding analysis establishes which translation directions are feasible at scale. The choice  
815 of API pairs for evaluation is motivated not only by data availability (Section C.1) but also by the  
816 *structural transformation complexity* each direction demands. We classify translation directions into  
817 four categories:

818 **Syntactic renaming** (e.g., CUDA to HIP): Near-1:1 API call mapping. `cudaMalloc` becomes  
819 `hipMalloc`; kernel launch syntax (`<<grid, block>>`) and thread-index arithmetic (`threadIdx.x`,  
820 `blockIdx.x`) are identical in HIP. This category primarily tests API surface adaptation rather than  
821 parallel reasoning.

```
822 // CUDA                                // HIP (syntactic rename)
823 #include <cuda.h>                        #include <hip/hip_runtime.h>
824 cudaMalloc(&d_a, size);                 hipMalloc(&d_a, size);
825 kernel<<<grid, block>>>(d_a);           kernel<<<grid, block>>>(d_a);
826 cudaMemcpy(h_a, d_a, ...);              hipMemcpy(h_a, d_a, ...);
```

Listing 2: CUDA to HIP: syntactic renaming.

829 **Paradigm translation** (e.g., CUDA to OpenMP): SPMD to fork-join model transformation. Ex-  
830 plicit GPU thread indexing must be converted to loop-based parallelism with OpenMP direc-  
831 tives. Shared memory and synchronization primitives (`__shared__`, `__syncthreads()`) must  
832 be replaced with stack-allocated arrays and implicit barrier semantics. Memory management  
833 (`cudaMalloc/cudaMemcpy/cudaFree`) is entirely eliminated.

```
834 // CUDA                                // OpenMP
835 __global__ void kern(float* a, int n) { void kern(float* a, int n) {
836     int i = blockIdx.x * blockDim.x    #pragma omp parallel for
837     + threadIdx.x;                     for (int i = 0; i < n; i++) {
838     if (i < n) a[i] = a[i] * 2.0f;      a[i] = a[i] * 2.0f;
839 }                                       }
840 }                                       }
841 }
```




Figure 6: Kernel-level API co-occurrence matrix. Each cell counts the number of individual kernels with implementations in both the row and column APIs. The CUDA–HIP pair leads with 633 shared kernels, followed by CUDA–SYCL (616), HIP–SYCL (615), and CUDA–OpenMP (472). OpenCL has dramatically lower kernel-level coverage: only 27 kernels share implementations with CUDA, and 24 with OpenMP. OpenMP Target provides 106 shared kernels with each of CUDA, HIP, SYCL, and OpenMP, concentrated entirely in RAJAperf. Sequential reference implementations exist for 106 kernels.




Figure 7: Kernel-level API co-occurrence network. Edges shown for  $\geq 20$  kernel pairs; edge width proportional to count. The tight CUDA/HIP/SYCL/OpenMP cluster reflects HeCBench and RAJAPerf’s uniform multi-API coverage. OpenCL appears as a peripheral node with weak connections (27 CUDA–OpenCL kernel pairs, just above the 20-kernel threshold). RAJA, stdpar, and Thrust appear as isolated nodes with fewer than 20 kernel-level pairs.




Figure 8: Repository-level vs. kernel-level translation pair counts. Left bars show the number of repositories containing both APIs; right bars show the number of independent kernel-level translation pairs. The multipliers (79 $\times$ –308 $\times$ ) show why PARBENCH counts kernels rather than repositories when estimating evaluation opportunity.

## Listing 3: CUDA to OpenMP: paradigm translation.

843 **Architecture split** (e.g., CUDA to OpenCL): CUDA’s single-source compilation model (host and  
 844 device code in one .cu file) must be split into separate host code (OpenCL runtime API calls for  
 845 platform/device/context/queue management) and kernel source (.cl files with OpenCL C syntax).  
 846 The memory model changes from CUDA’s device-pointer allocation (cudaMalloc/cudaMemcpy)  
 847 to OpenCL’s buffer objects (clCreateBuffer/clEnqueueWriteBuffer). Kernel launch syntax  
 848 transforms from `<<grid, block>>` to `clEnqueueNDRangeKernel` with explicit work-group sizing.

849 **Directive insertion** (e.g., sequential C to OpenMP): Adding parallelism to sequential code by  
 850 identifying parallelizable loops and inserting appropriate `#pragma omp` directives with correct data-  
 851 sharing clauses (`shared`, `private`, `reduction`). This is the inverse of the paradigm translation  
 852 direction and tests whether the LLM can *discover* parallelism rather than *translate* it.

853 PARBENCH targets the middle of this difficulty spectrum: paradigm translation and architecture split  
 854 (CUDA/OpenMP/OpenCL cross-translations). Syntactic renaming (CUDA to HIP) is too easy to be  
 855 informative, while directive insertion requires sequential baselines outside the current corpus. The six  
 856 directions among CUDA, OpenMP, and OpenCL span paradigm translation (CUDA $\leftrightarrow$ OpenMP) and  
 857 architecture split (CUDA $\leftrightarrow$ OpenCL, OpenMP $\leftrightarrow$ OpenCL), maximizing the semantic challenge while  
 858 remaining tractable for automated verification.

859 Table 6 summarizes the programming-model differences that make these directions structurally  
 860 different translation tasks.

Table 6: Programming model characteristics of the three primary translation APIs. These differences define the structural transformation challenges that LLM-based translation must address.

| Dimension                  | CUDA                                                                  | OpenMP                                                       | OpenCL                                                                        |
|----------------------------|-----------------------------------------------------------------------|--------------------------------------------------------------|-------------------------------------------------------------------------------|
| Execution model            | GPU SPMD kernels launched from host                                   | CPU fork-join threads via pragmas                            | Split host/device model with runtime kernel compilation                       |
| Thread/work-item indexing  | Explicit built-ins ( <code>threadIdx</code> , <code>blockIdx</code> ) | Loop indices and scheduling clauses                          | Explicit built-ins ( <code>get_global_id</code> , <code>get_local_id</code> ) |
| Memory management          | Explicit device allocation/free                                       | Shared host memory by default                                | Buffer objects plus explicit enqueue/write/read                               |
| Compilation model          | Single-source ahead-of-time GPU compilation                           | Host compiler with pragma lowering                           | Host code plus JIT-compiled kernel source                                     |
| File structure tendency    | Often single-source or host+kernel CUDA files                         | Usually host-language source only                            | Separate host source and .cl kernel file                                      |
| Typical translation burden | Remove explicit GPU constructs or introduce them                      | Introduce or remove compiler directives and loop parallelism | Split or merge host/device logic; satisfy OpenCL runtime constraints          |

861 **C.3 HIP and SYCL Exclusion Rationale**

862 Despite the large kernel-level co-occurrence counts for CUDA–HIP (633) and CUDA–SYCL (616),  
 863 these pairs were excluded from the primary evaluation for distinct reasons.

864 **HIP exclusion:** CUDA-to-HIP translation is predominantly a syntactic renaming task. The AMD HIP  
 865 API was explicitly designed as a drop-in replacement for CUDA, preserving the SPMD programming  
 866 model, kernel launch syntax, and memory management API. Automated tools such as `hipify-perl`  
 867 and `hipify-clang` achieve near-perfect conversion rates. Evaluating LLMs on CUDA-to-HIP  
 868 translation would measure API name lookup ability rather than parallel programming reasoning—the  
 869 core capability PARBENCH is designed to assess.

870 **SYCL exclusion:** While SYCL introduces a genuinely different programming model (single-source  
 871 C++ with lambda-based kernel dispatch and buffer/accessor memory management), it preserves the  
 872 GPU execution model. CUDA-to-SYCL translation is more complex than CUDA-to-HIP but remains  
 873 within the same architectural paradigm (GPU offload). Furthermore, SYCL benchmark coverage is  
 874 concentrated predominantly in HeCBench (approximately 80% of kernel-level instances), limiting  
 875 the diversity of evaluation instances. The CUDA-to-OpenMP direction provides a more fundamental

876 paradigm crossing (GPU SPMD to CPU fork-join) with better benchmark diversity across Rodinia,  
877 HeCBench, XSBench, and RSBench.

#### 878 C.4 OpenACC Exclusion

879 OpenACC was excluded despite appearing in 17 benchmarks in the repository-level survey (Figure 5),  
880 of which 12 also provide both CUDA and OpenMP implementations.

881 Three factors motivated this decision:

- 882 1. **Low kernel-level coverage:** At the kernel level, only 22 kernels provide OpenACC im-  
883 plementations across the surveyed repositories (Figure 6)—fewer than OpenCL (27) and  
884 substantially fewer than the primary APIs (CUDA: 656, OpenMP: 472). While repository-  
885 level co-occurrence with CUDA and OpenMP is high, the kernel-level material is insufficient  
886 for the multi-direction evaluation design that PARBENCH requires.
- 887 2. **Paradigm overlap with OpenMP target:** OpenACC’s directive-based model (`#pragma`  
888 `acc parallel loop`) occupies a similar conceptual niche to OpenMP’s target offload  
889 directives (`#pragma omp target teams distribute`). Including both would provide  
890 diminishing returns in programming-model diversity. OpenCL, despite comparable kernel  
891 counts, is retained because it exercises a qualitatively distinct programming model (explicit  
892 host/kernel separation and JIT compilation) not covered by any other included API.
- 893 3. **Compiler availability:** OpenACC compilation requires the NVIDIA HPC SDK  
894 (`nvc/nvc++`) or GCC with `-fopenacc`. This is a less universally available toolchain  
895 than CUDA (`nvcc`) or OpenMP (any modern C/C++ compiler), introducing a confounding  
896 variable between compiler availability and LLM translation capability.

#### 897 C.5 OpenMP Target as Case Study

898 OpenMP target offload (`#pragma omp target`) occupies a middle ground between CPU OpenMP  
899 and GPU-native APIs. It uses the OpenMP directive model but targets GPU execution, requiring  
900 the programmer to specify data mapping (`#pragma omp target data map(...)`) and work  
901 distribution (`#pragma omp teams distribute parallel for`) explicitly.

902 PARBENCH evaluates OpenMP target as a case study rather than a primary direction for three reasons:

- 903 1. **Compiler requirement:** OpenMP target compilation for NVIDIA GPUs requires the  
904 NVIDIA HPC compiler (`nvc/nvc++`, part of the NVIDIA HPC SDK 24.3). Default GCC  
905 and Clang installations support only CPU-threaded OpenMP; GPU offloading requires  
906 custom builds with target-offload support enabled. This introduces a confounding variable:  
907 a failed build may indicate compiler unavailability rather than translation quality.
- 908 2. **Limited benchmark coverage:** The kernel-level survey identifies 106 kernels with both  
909 OpenMP target and CUDA implementations (Figure 6), concentrated entirely in RAJAPerf.  
910 Within PARBENCH’s curated corpus, OpenMP target implementations are available for  
911 12 kernels: 10 from HeCBench, 1 from XSBench, and 1 from RSBench. Rodinia does not  
912 provide OpenMP target variants.
- 913 3. **Compilation model difference:** OpenMP target generates GPU offload code through the  
914 compiler, while CPU OpenMP generates threaded CPU code. A translation from CUDA  
915 to OpenMP target preserves GPU execution semantics but changes the syntax entirely; a  
916 translation from CUDA to CPU OpenMP changes both the execution model and the syntax.  
917 These are qualitatively different evaluation targets that should not be conflated in aggregate  
918 statistics.

919 The case study results for OpenMP target are reported separately in the direction analysis (Appendix  
920 Table 11) to enable direct comparison with the primary CUDA/OpenMP/OpenCL directions without  
921 confounding the aggregate pass rates.

## D Benchmark Kernel Survey

This appendix documents the systematic survey and multi-stage selection process that produced PARBENCH’s evaluation corpus. The main paper (Section 3) summarizes this process; here we provide the complete data, selection criteria, and exclusion rationale.

### D.1 Full Repository Inventory

The benchmark landscape survey covered 71 open-source HPC repositories (Appendix C.1). After applying the inclusion criteria described in Section D.2, 35 repositories remained as candidates for kernel extraction, spanning five categories of HPC benchmark software (Figure 9).




Figure 9: Distribution of benchmark types across the 35 repositories that met inclusion criteria. Suites (13, 37%) and mini-applications (12, 34%) dominate, accounting for 71% of included repositories. Proxy applications (4, 11%), libraries (4, 11%), and full applications (2, 6%) complete the distribution. Suites and mini-applications yield the richest kernel-level material for translation evaluation: they contain multiple independent computational kernels with self-checking verification.

API coverage varies widely across repositories: some implement a single API (e.g., SPEC OMP, LonestarGPU, STREAM), while others provide implementations across 10+ APIs (BabelStream and miniBUDE each with 12).

Table 7: Quantitative characterization of the evaluation corpus (35 kernels, 96 specs across 5 suites). SLoC counts physical source lines (non-blank, non-comment) in each kernel’s CUDA prompt payload, the code provided to the LLM as translation input. Multi-File gives specs requiring translation of >1 source file. Std. shows the dominant language standard per suite.

| Suite              | SLoC Range      | Med.       | Multi-File         | Std.  |
|--------------------|-----------------|------------|--------------------|-------|
| Rodinia            | 195–3,304       | 334        | 21/60 (35%)        | C++14 |
| HeCBench (curated) | 80–235          | 180        | 0/25 (0%)          | C++17 |
| XSbench            | 1,390           | –          | 2/4 (50%)          | C11   |
| RSBench            | 1,016           | –          | 1/4 (25%)          | C11   |
| mixbench           | 312             | –          | 0/3 (0%)           | C++11 |
| <b>Total</b>       | <b>80–3,304</b> | <b>271</b> | <b>24/96 (25%)</b> |       |

The API breadth distribution across the full 71-repository survey reveals the selection opportunity (Figure 10).

Taken together, these distributions show that kernel-level translation material is concentrated in a small number of large suites with broad API coverage—a Pareto distribution that guided suite selection. HeCBench and Rodinia were selected as the primary contributors due to their kernel richness and three-API coverage; XSbench, RSBench, and mixbench were added to extend domain coverage to nuclear physics simulation and GPU roofline characterization.




Figure 10: Distribution of API breadth across all 71 surveyed repositories. The mode is 3 APIs (18 repositories); a long tail extends to 12 APIs (BabelStream and miniBUDE). 24 repositories support only 1–2 APIs and are unsuitable for cross-API translation evaluation. Repositories supporting  $\geq 3$  APIs constitute the candidate pool for multi-direction evaluation.

## 940 D.2 Quality Tier Classification

941 Of the 71 surveyed repositories, 36 were excluded during initial screening: 31 lacked multi-API HPC  
 942 kernel content or fell outside the CUDA/OpenMP/OpenCL scope, and 5 had inaccessible downloads,  
 943 missing documentation, or duplicate content. The remaining 35 repositories were classified into two  
 944 quality tiers based on three criteria: build documentation, verification capability, and maintenance  
 945 status.

946 **Tier A criteria** (all three required):

- 947 • Documented build process (Makefile, CMake, or equivalent with clear instructions)
- 948 • Automated verification: self-checking output (print PASS/FAIL, compute checksum, com-  
 949 pare against reference) or reference output file comparison
- 950 • Active maintenance: commits within 2 years of the survey date, or stable release with no  
 951 known build failures on modern toolchains

952 **Tier B criteria** (any one triggers downgrade):

- 953 • Partial verification only (e.g., visual output inspection, manual comparison)
- 954 • Limited API coverage (fewer than three of the target APIs, reducing cross-translation utility)
- 955 • Unmaintained but still buildable (no commits in 2+ years, but Makefiles still produce  
 956 working binaries)

957 Tier B repositories were not excluded from the survey but were deprioritized in kernel selection.  
 958 Their kernels were considered only when no Tier A alternative existed for a given computational  
 959 domain.

## 960 D.3 Candidate Ranking and HeCBench Selection Funnel

961 The selection of individual kernels for PARBENCH’s evaluation corpus required balancing two  
 962 competing objectives: *kernel richness* (maximizing the number of independent evaluation instances)  
 963 and *API breadth* (maximizing the number of translation directions per kernel). Figure 11 plots these  
 964 two dimensions for all surveyed repositories.

965 Figure 12 presents the top candidates ranked by their combined kernel-count and API-breadth scores,  
 966 revealing the fundamental tradeoff in corpus construction.




Figure 11: Benchmark kernel richness (y-axis, capped at 100 for readability) vs. API breadth (x-axis, number of APIs supported). Bubble size proportional to total kernel count; color indicates benchmark type. HeCBench (522 kernels, 6 APIs<sup>3</sup>) dominates the upper region. RAJAPerf (80 kernels, 5 APIs) is the second-largest. CloverLeaf (23 kernels, 11 APIs) and BabelStream (5 kernels, 12 APIs) occupy the high-API-breadth region but with far fewer kernels. The optimal candidates for corpus construction lie in the high-kernel, moderate-API quadrant (upper-center).




Figure 12: Top multi-API repository candidates ranked by kernel count (blue bars, scaled for readability; true counts annotated above) and API breadth (orange bars). HeCBench provides the largest kernel pool (522 kernels, 6 APIs spanning CUDA, HIP, SYCL, OpenMP, OpenMP-target offloading, and MPI). BabelStream and miniBUDE provide the broadest API coverage (12 APIs each) but with few kernels. RAJAPerf (80 kernels, 5 APIs) offers a balanced middle ground. Rodinia (22 kernels, 4 APIs including a community SYCL port) provides established, well-understood benchmarks with strong community recognition.

#### 967 D.4 HeCBench Selection Funnel

968 HeCBench’s 522 kernels were filtered through a five-stage funnel to identify the curated evaluation  
969 set:




Figure 13: HeCBench kernel selection funnel. Starting from 522 kernels spanning CUDA, HIP, SYCL, and OpenMP at commit 22785cdd, successive filters for API completeness (329 with all 4 variants), build infrastructure (327), self-checking output patterns (242), and complexity bounds yield a 60-kernel working set. Manual curation further narrows this to 10 evaluation kernels (25 specs; see Section 3).

- 970 • **Stage 1 (API completeness):** Of 522 total kernels in the HeCBench repository, 329 provide  
971 implementations in all 4 primary APIs (CUDA, HIP, SYCL, OpenMP).
- 972 • **Stage 2 (Build infrastructure and verification):** Of these 329 kernels, 2 lack Makefiles and  
973 were excluded. Of the remaining 327, only 242 produce deterministic, machine-parseable  
974 self-checking output—printing PASS/FAIL verdicts, computing checksums, or comparing  
975 numerical results against reference values—as determined by manual inspection of the  
976 CUDA source at commit 22785cdd. These 242 candidates proceed to complexity filtering.
- 977 • **Stage 3 (Complexity filter):** The 242 candidates were filtered for evaluation tractability.  
978 Kernels with more than 15 source files or fewer than 2 were excluded (the former are too  
979 complex for single-prompt translation; the latter are trivial boilerplate). Kernels requiring  
980 external input data files were also excluded to ensure self-contained evaluation.
- 981 • **Stage 4 (Domain diversity):** From the complexity-filtered set, 60 kernels were manually  
982 selected to ensure coverage across 10 computational categories: linear algebra, stencil  
983 computation, graph algorithms, machine learning, physics simulation, sorting, reduction,  
984 scan, image processing, and cryptography.
- 985 • **Stage 5 (Curation):** The 60-kernel working set was narrowed to a final curated set of 20  
986 kernels based on manual inspection of build reliability, output determinism, and domain  
987 representativeness. Of these, 10 kernels were included in the current PARBENCH evaluation  
988 corpus, with the remaining 10 reserved for future expansion.

989 The 10 curated kernels included in the current corpus are: convolution1d, floydwarshall,  
990 heat2d, iso2dfd, jacobi, md, nqueen, page-rank, scan, and stencil1d. Each provides a  
991 CUDA implementation paired with one or more OpenMP/OpenMP-target variants, totaling 25 specs.  
992 Two of these are KNOWN\_FAIL: stencil1d-omp\_target (build failure due to target-offloading  
993 compilation) and scan-omp\_target (output mismatch on the CPU offload target), leaving 23 in the  
994 active evaluation corpus.

995 This funnel illustrates a general challenge in constructing LLM code translation benchmarks: the  
 996 vast majority of available HPC kernels are unsuitable for automated evaluation due to missing API  
 997 variants, non-deterministic output, or excessive complexity. Only 3.8% of HeCBench’s 522 kernels  
 998 (20/522) survive the full selection pipeline, highlighting the gap between “available parallel code”  
 999 and “evaluable parallel code.”

## 1000 D.5 Rodinia Kernel Inventory

1001 Rodinia Che et al. [2009] provides the foundation of PARBENCH’s evaluation corpus. As one  
 1002 of the most widely cited heterogeneous computing benchmark suites, it offers well-understood  
 1003 computational characteristics across CUDA, OpenMP, and OpenCL. PARBENCH includes all 22  
 1004 Rodinia kernels, yielding 60 total specs; five kernels lack one or more API variants because the  
 1005 upstream repository does not provide implementations for all three APIs (six API variants absent in  
 1006 total, as one kernel—`huffman`—lacks both OpenMP and OpenCL). Of these 60 specs:

- 1007 • **53 PASS:** Each verified via conjunctive application of all its oracle strategies—every strategy  
 1008 must pass for the spec to pass. 48 of the 53 use exit-code-zero *and* stdout pattern matching.  
 1009 The remaining five employ stronger oracle methods: three (`hotspot3d`) augment the  
 1010 conjunctive check with numeric comparison against baseline output within a fixed tolerance,
- 1011 • **7 KNOWN\_FAIL:** Excluded from evaluation due to pre-existing platform failures unrelated  
 1012 to LLM translation quality: 2 specs use deprecated CUDA `texture<>` references removed  
 1013 in CUDA 12 (`kmeans-cuda`, `mummergpu-cuda`); 1 requires OpenGL (`hybridsort-cuda`);  
 1014 1 has CUDA API dependencies in its OpenMP variant (`mummergpu-omp`); 2 exhibit pre-  
 1015 existing OpenCL segmentation faults (`nn-openc1`, `kmeans-openc1`); and 1 has a silent  
 1016 `clEnqueueReadBuffer` error masked by exit code 0 (`backprop-openc1`).

## 1017 D.6 XSBench, RSBench, and mixbench Profiles

1018 Three additional benchmark suites complement Rodinia and HeCBench in the evaluation corpus:

1019 **XSBench** Tramm et al. [2014] (4 specs: CUDA, OpenMP, OpenCL, OpenMP Target): A Monte  
 1020 Carlo neutron cross-section lookup proxy application representative of nuclear reactor simulation  
 1021 workloads. XSBench implements a unionized energy grid algorithm; the OpenMP variant runs in  
 1022 history-based mode (checksum 941,535) while the CUDA, OpenCL, and OpenMP Target variants run  
 1023 in event-based mode (checksum 945,990). Both checksums are self-verified as valid by XSBench’s  
 1024 built-in verification. This mode asymmetry reflects algorithmic differences between CPU and GPU  
 1025 scheduling strategies rather than correctness issues. All 4 API variants pass baseline verification.

1026 **RSBench** Tramm and Siegel [2014] (4 specs: CUDA, OpenMP, OpenCL, OpenMP Target): A proxy  
 1027 application for the multipole method of computing neutron cross sections on-the-fly, representing  
 1028 a complementary algorithm to XSBench’s unionized energy grid approach. Like XSBench, the  
 1029 OpenMP variant defaults to history-based mode (checksum 879,693) while the CUDA, OpenCL, and  
 1030 OpenMP Target variants run in event-based mode (checksum 880,018), with both checksums self-  
 1031 verified as valid. Including both suites enables controlled comparison within the same computational  
 1032 domain (Monte Carlo neutron transport) at different algorithmic complexity levels. All 4 API variants  
 1033 pass baseline verification.

1034 **mixbench** Konstantinidis and Cotronis [2017] (3 specs: CUDA, OpenMP, OpenCL): A micro-  
 1035 benchmark that sweeps a range of compute-to-memory operation ratios, mapping the practical  
 1036 roofline boundary of a compute device. Its parameterized kernel structure tests whether LLM  
 1037 translations preserve arithmetic intensity across API boundaries. All 3 API variants pass baseline  
 1038 verification.

## 1039 D.7 Exclusion Log and Verification Methods

1040 Figure 14 groups the surveyed benchmarks by their verification approach. The dominant clusters—  
 1041 checksum validation and reference output comparison—contain the majority of Tier A benchmarks  
 1042 (Section D.2), while benchmarks relying solely on visual inspection or manual comparison were  
 1043 excluded or downgraded to Tier B.




Figure 14: Benchmarks grouped by verification method. Large nodes represent verification approaches; benchmark nodes are connected to their method(s). Benchmarks without automated verification were excluded or downgraded to Tier B (Section D.2). The checksum and reference output clusters contain the majority of Tier A benchmarks.

## Repository-Level Exclusions

Of the 71 surveyed repositories, 36 were excluded in two stages. First, 31 were removed during initial screening (fewer than three API implementations, no HPC kernel content, or outside the CUDA/OpenMP/OpenCL scope), leaving 40 candidates. Of these, five were excluded during detailed evaluation:

- **Download failures** (2): Repository URLs were dead or access-restricted at the time of survey.
- **Insufficient documentation** (2): No build instructions, no Makefiles, or build system required proprietary dependencies not available on the evaluation platform.
- **Duplicate content** (1): One repository was a fork of an already-surveyed suite with no additional API variants.

## Kernel-Level Exclusions

Within the selected suites, individual kernels were excluded for the following reasons:

- **Insufficient API coverage:** Kernels lacking the API variants needed for multi-direction translation evaluation. The completeness threshold is suite-specific: HeCBench required all four primary APIs (Stage 1 of the selection funnel, Section D.3); Rodinia required at least two of CUDA, OpenMP, and OpenCL.

- **No self-checking output:** Kernels that produce graphical output, require manual inspection, or have no deterministic expected output.
- **Excessive complexity:** Kernels requiring more than 15 source files (the HeCBench selection filter), external data files not included in the repository (e.g., Rodinia’s `leukocyte` requires a separately downloaded video dataset and 102 external library files), or runtime dependencies (e.g., OpenGL for Rodinia’s `hybridsort`, MPI multi-node) not available on the single-GPU evaluation platform.
- **Non-deterministic output:** Kernels whose output depends on execution order, random seeds without fixed initialization, or floating-point non-associativity that prevents consistent cross-API baseline matching. (Retained kernels with borderline FP divergence—such as `cf_d`, `hotspot`, and `myocyte`—use relaxed stdout-pattern oracles rather than hash-based verification; see Section D.5.)

For kernels retained in the corpus but failing baseline verification (KNOWN\_FAIL specs), the specific failure reason and build output are recorded in each spec’s `baseline_results` field. For kernels excluded before spec creation (e.g., Rodinia’s `leukocyte`), the exclusion is documented in the survey data and selection funnel (Sections D.3–D.5).

## E Augmentation Protocol and Extended Evaluation Results

This appendix collects the augmentation protocol details that would interrupt the main text, then groups the result-heavy supporting evidence in the same order as the main paper’s empirical claims. The first subsection defines the L0–L4 levels used throughout the robustness analysis. The next subsection documents transform frequency and baseline validation limits. The remaining subsections cover a representative multi-file failure case and the extended result tables that support the  $\text{pass}@k$ , augmentation, and kernel-difficulty claims in Section 5.

### E.1 Augmentation Levels

Table 8 gives the exact L0–L4 augmentation definitions used in the robustness protocol.

Table 8: Augmentation level definitions.

| Level | Transform selection | Candidate fraction | Description                     |
|-------|---------------------|--------------------|---------------------------------|
| L0    | None                | 0%                 | Unmodified source code baseline |
| L1    | 1 transform         | 1 site             | Minimal source perturbation     |
| L2    | 33% of transforms   | 33% of candidates  | Moderate perturbation           |
| L3    | 66% of transforms   | 66% of candidates  | Heavy perturbation              |
| L4    | All transforms      | 100% of candidates | Maximum configured perturbation |

### E.2 Augmentation Transform Frequency And Baseline Validation

PARBENCH’s code augmentation pipeline applies six AST-level transforms at four intensity levels (L1–L4). The transforms do not apply uniformly: some kernels have many candidate sites for a given transform, while others have few or none. Figure 15 presents the per-kernel, per-transform frequency heatmap.

Aggregate transform frequencies across the corpus reveal a clear dominance hierarchy. Below we report the number of Rodinia specs (out of 60) to which each transform applies at L4 (full intensity), followed by the total application count across all 240 augmentation tasks ( $60 \text{ specs} \times 4 \text{ levels}$ ):

- **SwapCondition** (59/60 specs at L4; 162 applications across all levels): The most broadly applicable transform. Swaps operands of comparison operators (e.g.,  $a < b \rightarrow b > a$ ), with guards against expressions containing side effects. Its near-universal coverage reflects the ubiquity of comparison-based control flow in HPC kernels—boundary checks, convergence tests, and branch-based algorithm selection all contribute candidate sites.




Figure 15: Per-kernel transform application frequency across the Rodinia corpus. Rows represent the 22 Rodinia kernels; columns represent the six AST-level transforms. Cell values give the total number of times each transform was applied to files belonging to that kernel, aggregated across all augmentation levels and API variants. The distribution is highly non-uniform: multi-file kernels such as `myocyte` (16 source files) and `bptree` accumulate many `SwapCondition` applications, while smaller kernels (e.g., `pathfinder`, `hotspot`) show lower counts. `ChangeFunctionNames` applies to only two kernels (`huffman`, `streamcluster`).

- 1099 • **ArithmeticTransform** (45/60 specs at L4; 69 applications): Converts compound assignment  
1100 operators to their expanded forms and vice versa (e.g., `x += expr`  $\leftrightarrow$  `x = x + (expr)`).  
1101 Expansion applies to integer and floating-point scalar targets; folding back to compound  
1102 form is restricted to integer-typed scalars to avoid floating-point precedence issues. For-loop  
1103 incrementors are excluded to avoid disrupting iteration patterns. Less broadly applicable than  
1104 `SwapCondition` because it requires compound-assignment or assignment-with-matching-  
1105 LHS patterns that not all kernels exhibit.
- 1106 • **ChangeNames** (32/60 specs at L4; 55 applications): Consistently renames local variables  
1107 and function parameters using symbol-aware reference collection, with behavior-preserving  
1108 intent validated against the retained baseline subset. Coverage is bounded by the number of  
1109 renameable local declarations per kernel.
- 1110 • **TypedefExpansion** (3/60 specs at L4; 7 applications): Inlines `typedef` aliases to their  
1111 underlying types. Rare because most Rodinia kernels use primitive types directly rather than  
1112 defining type aliases.
- 1113 • **PointerArithmeticToArrayIndex** (3/60 specs at L4; 6 applications): Rewrites `*(ptr +`  
1114 `i)` to `ptr[i]` and vice versa, with precedence-aware handling of struct member access.  
1115 Limited by the prevalence of pointer arithmetic in the corpus—most kernels use array  
1116 indexing natively.
- 1117 • **ChangeFunctionNames** (2/60 specs at L4; 2 applications): Renames user-defined functions  
1118 and updates all call sites, skipping `main`, GPU kernel entry points, and identifiers appearing  
1119 in string literals. The rarest transform, reflecting the small number of user-defined functions  
1120 in single-kernel files.

The level-invariance property supports that augmentation introduces no new baseline failures on the retained Rodinia subset: all 53 non-`KNOWN_FAIL` Rodinia specs pass at every augmentation level L1–L4. (Note that the per-transform applicability fractions above, e.g., 59/60 for `SwapCondition`, are computed over all 60 Rodinia specs including the 7 `KNOWN_FAIL` specs, since augmentation is applied to source files regardless of baseline status; baseline validation uses only the 53 non-`KNOWN_FAIL` specs.) All transforms are designed to preserve behavior, supported by 15 unit tests covering each transform in isolation and by the end-to-end augmentation sweep across the full Rodinia corpus.

Across all five suites, 80 of 87 non-`KNOWN_FAIL` specs pass at all levels L1–L4, with the 7 failures confined to `omp_target` variants. Five of these (all `HeCBench` kernels) fail at Level 3 with `SwapCondition` as the sole active transform, directly implicating condition operand swapping (e.g.,  $a < b \rightarrow b > a$ ) in causing the GPU-offloading compiler to generate numerically divergent device code. The remaining two (`rsbench`, `xsbench`) fail only at Level 4 where both increased `SwapCondition` intensity and the addition of `ArithmeticTransform` coincide, making precise attribution ambiguous. Variable renaming is separately disabled for any file containing `#pragma omp` directives, so it plays no role in these failures. All 87 specs pass at L1–L2. Because these baseline failures are specific to the `omp_target` toolchain, the affected variants are excluded from robustness claims about `omp_target` baselines and still serve only as prompt inputs when translating to other targets.

### E.3 Case Study: Multi-File Translation Consistency Failure

Multi-file kernel structures expose direction-dependent translation failures that single-file kernels do not exhibit. We document these through a detailed case study of the `rodinia-lavamd` kernel, a molecular dynamics simulation with a multi-file build structure (kernel, wrapper, and supporting headers). `PARBENCH` evaluates `lavamd` in both full-program mode (CUDA/OMP targets, where the LLM rewrites kernel and host code) and kernel-only mode (OpenCL targets, where only the `.c1` compute kernel is translated), revealing distinct failure mechanisms in each.

#### Observation

Across all six standard translation directions, Qwen 3.5 fails every `lavamd` translation: 0 of 18 records (6 directions  $\times$  3 independent samples at temperature 0.7) produce a `PASS` outcome. The failure mode is not uniform: it varies systematically by translation direction rather than by stochastic sample, revealing that the failure mechanism is structurally determined by the source–target API pairing.

#### Per-Direction Error Taxonomy

Table 9: Translation error taxonomy for `lavamd` across all six standard directions (Qwen 3.5, L0, 3 samples each). Full-program directions (CUDA/OMP targets) fail deterministically at build time; kernel-only directions (OpenCL targets) fail at runtime. Five of six directions produce identical failure modes across all three samples.

| Direction   | Status ( $\times 3$ ) | Primary Failure Mode                                  |
|-------------|-----------------------|-------------------------------------------------------|
| OpenCL→CUDA | BUILD_FAIL            | Phantom include ( <code>timing.h</code> )             |
| OMP→CUDA    | BUILD_FAIL            | Phantom include ( <code>./util/timer/timer.h</code> ) |
| CUDA→OMP    | BUILD_FAIL            | Phantom include ( <code>./main.h</code> )             |
| CUDA→OpenCL | RUN_FAIL              | OpenCL kernel compilation failure                     |
| OMP→OpenCL  | RUN_FAIL              | OpenCL kernel compilation failure                     |
| OpenCL→OMP  | Mixed                 | 2 BUILD_FAIL + 1 VERIFY_FAIL                          |

Table 9 reveals two structural patterns corresponding to `PARBENCH`’s translation modes. First, three of four full-program directions (targeting CUDA or OMP, where the LLM rewrites both kernel and host code) fail uniformly at build time due to phantom include paths: the model generates headers (`timing.h`, `./main.h`, `./util/timer/timer.h`) that exist in the source program’s broader directory tree but are unavailable in the translated program’s build context. All three samples within each of these directions fail on the same phantom header, indicating a deterministic error rather than stochastic variation. The fourth full-program direction (OpenCL→OMP) exhibits mixed failure

1159 modes across its three samples—including a phantom include, an OpenMP barrier-nesting violation,  
 1160 and a verification failure where the program runs to completion (exit code 0) but produces only  
 1161 configuration output, no computation results. Second, the two kernel-only directions (targeting  
 1162 OpenCL, where only the .cl kernel file is translated while the host driver is untouched) pass host-  
 1163 code compilation but fail at runtime when `clBuildProgram` attempts to compile the translated  
 1164 kernel source (exit code 245 = `CL_BUILD_PROGRAM_FAILURE`). The runtime compilation errors  
 1165 include incorrect `__local` variable scoping, undeclared identifiers (e.g., `fp`, `NUMBER_THREADS`),  
 1166 and missing address space qualifiers on kernel pointer parameters—API-specific constraints that  
 1167 the model fails to satisfy. Cross-model comparison reveals that these barriers are not uniformly  
 1168 intractable: GPT-5.4 passes 5 of 18 records by inlining header definitions (avoiding phantom includes)  
 1169 and producing harness-passing OMP→OpenCL kernels, while GPT-5.3-codex passes 3 of 18 on the  
 1170 same kernel-only direction. The full-program CUDA-targeting directions, however, remain unsolved  
 1171 across all three models.

## 1172 Cross-Model Persistence

1173 The direction-dependent failure pattern persists across models. Only two of six directions yield  
 1174 any PASS outcomes across the three models: OMP→OpenCL (GPT-5.4 and GPT-5.3-codex each  
 1175 3/3 PASS, Qwen 3.5 0/3) and CUDA→OMP (GPT-5.4 2/3 PASS, Qwen 3.5 and GPT-5.3-codex  
 1176 0/3). The remaining four directions are universally failed across all models and samples (0/36 PASS).  
 1177 OpenCL→CUDA and OMP→CUDA fail at build time due to phantom includes (18/18 `BUILD_FAIL`),  
 1178 CUDA→OpenCL fails at runtime due to OpenCL kernel compilation errors (9/9 `RUN_FAIL`), and  
 1179 OpenCL→OMP fails with a mix of phantom includes and output mismatches. These asymmetries  
 1180 indicate that multi-file coordination difficulty is a property of the translation direction—not solely of  
 1181 the kernel complexity or the model.

## 1182 Implications

1183 This case study identifies three distinct error subcategories within the “multi-file translation consis-  
 1184 tency failure” class:

- 1185 1. **Phantom dependencies:** The model hallucinates include files (`timing.h`, `./main.h`) or  
 1186 helper functions (`checkCUDAError`, `setdevice`) that may exist in the model’s training  
 1187 data but are absent from the target program’s file structure. This is the dominant build-time  
 1188 failure mode for CUDA-targeting directions across all three models.
- 1189 2. **Cross-file identifier inconsistency:** In some translations (observed in GPT-5.3-codex  
 1190 and GPT-5.4 OpenCL→CUDA attempts), the kernel file uses the target API’s naming  
 1191 convention while the wrapper file retains the source API’s function name, producing linker  
 1192 errors (undefined reference to `‘kernel_gpu_cuda’`).
- 1193 3. **Runtime kernel compilation failure:** For OpenCL-targeting directions, the host C/C++  
 1194 code compile successfully, but the OpenCL kernel source, compiled at runtime by  
 1195 `clBuildProgram`, contains errors such as incorrect `__local` variable scoping and un-  
 1196 declared type aliases. This represents a disconnect between the C/C++ compiler’s view of  
 1197 the host code and the OpenCL runtime compiler’s requirements for device code, a failure  
 1198 mode invisible to static build checks.

1199 These findings suggest that multi-file translation failures are not uniformly distributed across direc-  
 1200 tions. Directions requiring the model to generate host-side boilerplate for a fundamentally different  
 1201 runtime (e.g., CUDA driver API, OpenCL `cl*` calls) are systematically harder than directions where  
 1202 the host code structure is similar between source and target. A hybrid pipeline combining LLM trans-  
 1203 lation for computational kernels with template-based scaffold generation for host infrastructure may  
 1204 address subcategories 1 and 2 but would not resolve the runtime compilation failures in subcategory 3,  
 1205 which require translated API-specific device-code constraints to satisfy the runtime compiler.

## 1206 E.4 Extended Results Details

1207 This section expands on the results presented in Section 5. Across all three models, we collected  
 1208 2,262 valid translation records: 626 for Qwen 3.5 (426 L0 + 200 augmentations, after excluding 82

records involving 9 KNOWN\_FAIL specs<sup>4</sup>), 822 for GPT-5.4 (426 L0 + 396 augmentations), and 814 for GPT-5.3-codex (426 L0 + 388 augmentations). KNOWN\_FAIL specs were pre-excluded from the GPT-5.4 and GPT-5.3-codex evaluation batches.

Table 10: Record-level diagnostic pass rates across three models and 2,262 valid translation records. Records include L0 samples and L1–L4 augmentation runs for L0-passing pairs, so this table mixes canonical L0 records with model-specific L0-conditional augmentation records. Wilson intervals are descriptive because records are clustered by task; inferential comparison uses the balanced 142 L0 pairs in Section 5.1.

| Model              | PASS | BUILD_FAIL | RUN_FAIL | VERIFY_FAIL | EXTRACT_FAIL | Total | Rate [95% Wilson CI] |
|--------------------|------|------------|----------|-------------|--------------|-------|----------------------|
| Qwen 3.5 397B-A17B | 230  | 245        | 121      | 29          | 1            | 626   | 36.7% [33.1%, 40.6%] |
| GPT-5.4            | 621  | 123        | 43       | 32          | 3            | 822   | 75.5% [72.5%, 78.4%] |
| GPT-5.3-codex      | 604  | 139        | 44       | 27          | 0            | 814   | 74.2% [71.1%, 77.1%] |

## Record-Level Diagnostic Summary.

Table 11: Pass rates by translation direction (L0 records, all three samples per task). Six standard directions form the core matrix; four OMP-target directions are case studies with fewer participating kernels ( $m$ ). Here  $n$  denotes L0 records per model for that direction. Wilson 95% CIs in brackets.

| Direction                 | Qwen 3.5             | GPT-5.4              | GPT-5.3-codex        | $n$ |
|---------------------------|----------------------|----------------------|----------------------|-----|
| CUDA→OMP                  | 40.3% [29.7%, 51.8%] | 83.3% [73.1%, 90.2%] | 76.4% [65.4%, 84.7%] | 72  |
| OMP→OpenCL                | 33.3% [22.0%, 47.0%] | 72.5% [59.1%, 82.9%] | 82.4% [69.8%, 90.4%] | 51  |
| OMP→CUDA                  | 25.0% [16.4%, 36.1%] | 55.6% [44.1%, 66.5%] | 55.6% [44.1%, 66.5%] | 72  |
| OpenCL→OMP                | 9.8% [4.3%, 21.0%]   | 41.2% [28.8%, 54.8%] | 39.2% [27.0%, 52.9%] | 51  |
| CUDA→OpenCL               | 7.0% [2.8%, 16.7%]   | 59.7% [46.7%, 71.4%] | 57.9% [45.0%, 69.8%] | 57  |
| OpenCL→CUDA               | 0.0% [0.0%, 6.3%]    | 19.3% [11.1%, 31.3%] | 19.3% [11.1%, 31.3%] | 57  |
| OMP-target→OMP ( $m=3$ )  | 100% [70.1%, 100%]   | 100% [70.1%, 100%]   | 100% [70.1%, 100%]   | 9   |
| OMP-target→CUDA ( $m=8$ ) | 66.7% [46.7%, 82.0%] | 95.8% [79.8%, 99.3%] | 100% [86.2%, 100%]   | 24  |
| OMP→OMP-target ( $m=3$ )  | 44.4% [18.9%, 73.3%] | 100% [70.1%, 100%]   | 100% [70.1%, 100%]   | 9   |
| CUDA→OMP-target ( $m=8$ ) | 0.0% [0.0%, 13.8%]   | 95.8% [79.8%, 99.3%] | 100% [86.2%, 100%]   | 24  |

## Direction-Level Pass Rates.

Table 12: pass@ $k$  results for Qwen 3.5 (L0 only, temperature 0.7, three independent samples per task, 142 unique source-target pairs). CIs are normal-approximation intervals on the macro-averaged Chen et al. estimator (distinct from the Wilson binomial CIs used for per-record pass rates in Table 10)

| Metric            | Value [95% CI]       |
|-------------------|----------------------|
| pass@1            | 23.9% [17.9%, 30.0%] |
| pass@3            | 35.2% [27.3%, 43.1%] |
| Always pass (3/3) | 19 tasks (13.4%)     |
| Noisy (1–2/3)     | 31 tasks (21.8%)     |
| Hard fail (0/3)   | 92 tasks (64.8%)     |

**pass@ $k$  Analysis.** The 36.7% per-record rate (Table 10) treats each sample and augmentation level independently. The pass@ $k$  analysis (Table 12) restricts to the 142 L0 tasks and macro-averages over tasks to estimate single-sample and three-sample success probabilities, with each task generating three independent samples at temperature 0.7. The gap between pass@1 (23.9%) and pass@3 (35.2%) is 11.3%: only 31 of 142 tasks show within-task variability (1 or 2 of 3 samples pass); none of the 92 hard failures (0/3) are recovered by resampling. This confirms that the majority of failures reflect systematic translation challenges—structural mismatches between parallel programming models—rather than stochastic variation, validating PARBENCH’s design as a benchmark that surfaces genuine capability gaps.

<sup>4</sup>The Qwen 3.5 total is 626 rather than 630 because four rodinia-backprop-openc1 L1–L4 augmentation records are also excluded via the KNOWN\_FAIL policy (the OpenCL baseline is silently broken).

Table 13: Pass@ $k$  estimates for the shared L0 task set (142 evaluation-eligible direction–kernel pairs, three independent samples per pair). Hard Fail = 0/3 samples pass; Always Pass = 3/3; Noisy = 1–2/3.

| Model         | Pairs | pass@1 | pass@3 | Hard Fail      | Noisy          | Always Pass    |
|---------------|-------|--------|--------|----------------|----------------|----------------|
| Qwen 3.5      | 142   | 23.9%  | 35.2%  | 92/142 (64.8%) | 31/142 (21.8%) | 19/142 (13.4%) |
| GPT-5.4       | 142   | 62.7%  | 69.7%  | 43/142 (30.3%) | 23/142 (16.2%) | 76/142 (53.5%) |
| GPT-5.3-codex | 142   | 62.7%  | 68.3%  | 45/142 (31.7%) | 16/142 (11.3%) | 81/142 (57.0%) |

1223 **Cross-Model pass@ $k$  Summary.** Table 13 extends the single-model view above to all three model  
1224 campaigns on the shared 142-task L0 set.

1225 The cross-model summary in Table 13 shows that GPT-5.4 and GPT-5.3-codex share the same pass@1  
1226 while differing slightly in consistency: GPT-5.3-codex has more always-pass pairs and fewer noisy  
1227 pairs, whereas Qwen 3.5 concentrates much more mass in hard failures. Across all three models,  
1228 most failed L0 tasks remain unrecovered by resampling.

1229 **Representative Per-Kernel Pass Rates.** Table 14 highlights the easiest and hardest kernels for  
Qwen 3.5.

Table 14: Representative per-kernel pass rates for Qwen 3.5 across all directions and augmentation levels (626 records, 31 evaluated kernels—4 curated kernels have no evaluable translation pairs). nqueen is tied with page-rank at 60.0%.

| Suite                                                         | Kernel        | Total | PASS | Fail | Rate  |
|---------------------------------------------------------------|---------------|-------|------|------|-------|
| <i>Easiest kernels</i>                                        |               |       |      |      |       |
| HeCBench                                                      | iso2dfd       | 38    | 32   | 6    | 84.2% |
| HeCBench                                                      | floydwarshall | 38    | 29   | 9    | 76.3% |
| HeCBench                                                      | heat2d        | 38    | 29   | 9    | 76.3% |
| HeCBench                                                      | stencil1d     | 14    | 9    | 5    | 64.3% |
| HeCBench                                                      | page-rank     | 10    | 6    | 4    | 60.0% |
| <i>Hardest kernels (0% pass rate, <math>n \geq 18</math>)</i> |               |       |      |      |       |
| Rodinia                                                       | heartwall     | 18    | 0    | 18   | 0.0%  |
| Rodinia                                                       | myocyte       | 18    | 0    | 18   | 0.0%  |
| Rodinia                                                       | bptree        | 18    | 0    | 18   | 0.0%  |
| RSBench                                                       | rsbench       | 18    | 0    | 18   | 0.0%  |
| Rodinia                                                       | lavamd        | 18    | 0    | 18   | 0.0%  |

1230

Table 15: Augmentation pass rates on the common 12-kernel CUDA-to-OpenMP subset present at all five levels. These 12 kernels were selected by the L0-conditional filter ( $\geq 1$  of 3 L0 samples passes), so L0 rates exceed the overall CUDA-to-OpenMP rate of 40.3%.

| Level | Pass / Total | Rate  | Kernels with any pass |
|-------|--------------|-------|-----------------------|
| L0    | 29/36        | 80.6% | 12/12 (100%)          |
| L1    | 9/12         | 75.0% | 9/12 (75%)            |
| L2    | 10/12        | 83.3% | 10/12 (83%)           |
| L3    | 9/12         | 75.0% | 9/12 (75%)            |
| L4    | 10/12        | 83.3% | 10/12 (83%)           |

1231 **Augmentation Pass Rates.** If baseline success were driven primarily by surface-form memorization,  
1232 augmented source variants should degrade pass rates. On a balanced 12-kernel CUDA-to-OpenMP  
1233 subset present at all five augmentation levels (Table 15), pass rates remain between 75% and 83% at  
1234 L1–L4. The L0 rate on this subset (80.6%) exceeds the overall CUDA-to-OpenMP L0 rate (40.3%)  
1235 because the L0-conditional filter selects kernels that already pass at L0, creating a survivorship  
1236 bias toward easier kernels. We considered Cochran–Armitage trend tests, but two confounds limit  
1237 their interpretability on this subset: the 3:1 ratio of L0 to L1–L4 sample sizes, and the survivorship  
1238 selection inherent in L0-conditional filtering. The observed stability is compatible with robustness to

Table 16: Declared-oracle pass rates across augmentation levels. L0 covers all 142 evaluation-eligible direction–kernel pairs ( $n = 426$  L0 records, three samples each). L1–L4 use one augmentation record per qualifying pair on model-specific L0-conditional subsets: pairs where at least one of three L0 samples passes (Qwen 3.5: 50 pairs, GPT-5.4: 99 pairs, GPT-5.3-codex: 97 pairs).

| Level | Qwen 3.5<br>(all dirs) | Qwen 3.5<br>(C→OMP) | GPT-5.4<br>(all dirs) | GPT-5.4<br>(C→OMP) | GPT-5.3-codex<br>(all dirs) | GPT-5.3-codex<br>(C→OMP) |
|-------|------------------------|---------------------|-----------------------|--------------------|-----------------------------|--------------------------|
| L0    | 23.9% ( $n=426$ )      | 40.3% ( $n=72$ )    | 62.7% ( $n=426$ )     | 83.3% ( $n=72$ )   | 62.7% ( $n=426$ )           | 76.4% ( $n=72$ )         |
| L1    | 74.0% ( $n=50$ )       | 75.0% ( $n=12$ )    | 88.9% ( $n=99$ )      | 90.9% ( $n=22$ )   | 86.6% ( $n=97$ )            | 85.0% ( $n=20$ )         |
| L2    | 64.0% ( $n=50$ )       | 83.3% ( $n=12$ )    | 90.9% ( $n=99$ )      | 95.5% ( $n=22$ )   | 88.7% ( $n=97$ )            | 85.0% ( $n=20$ )         |
| L3    | 62.0% ( $n=50$ )       | 75.0% ( $n=12$ )    | 86.9% ( $n=99$ )      | 81.8% ( $n=22$ )   | 86.6% ( $n=97$ )            | 85.0% ( $n=20$ )         |
| L4    | 56.0% ( $n=50$ )       | 83.3% ( $n=12$ )    | 90.9% ( $n=99$ )      | 90.9% ( $n=22$ )   | 85.6% ( $n=97$ )            | 90.0% ( $n=20$ )         |




Figure 16: L0-conditional augmentation outcomes across all three models for the CUDA-to-OpenMP direction. L0 uses first-sample results across all kernels; L1–L4 are restricted to each model’s qualifying subset. Error bars show Wilson 95% intervals and are descriptive because the subset differs by model.

1239 the tested surface-form perturbations, though it does not rule out memorization at higher levels of  
1240 abstraction (e.g., algorithmic templates).

1241 **All-Model Augmentation Summary.** Table 16 and Figures 16–17 extend the balanced-subset view  
1242 to all three model campaigns. These summaries remain descriptive because the L0-conditional filter  
1243 changes the denominator across models and levels.

1244 Table 16 shows the broader pattern behind the balanced 12-kernel subset: GPT-5.4 and GPT-5.3-codex  
1245 stay above 85% on their L0-conditional subsets across L1–L4, while Qwen 3.5 declines from L1 to  
1246 L4. These rates should still be read as surface-form robustness on the qualifying subset rather than as  
1247 a definitive memorization test.

1248 **Direction Asymmetry And Test Summary.** Table 17 consolidates the paired direction-asymmetry  
1249 tests and the descriptive augmentation summary used in the main-text interpretation.




Figure 17: Per-kernel L0/L1–L4 outcomes for the CUDA-to-OpenMP subset under Qwen 3.5. Cell color indicates whether the corresponding analysis record receives declared-oracle PASS or a failure outcome at that level.

Table 17: Statistical summary (Qwen 3.5 canonical evaluation). All confidence intervals are Wilson 95% intervals. McNemar tests in the direction asymmetry family use Bonferroni correction ( $\alpha = 0.05/5 = 0.01$ ). Augmentation results are summarized descriptively because the L0-conditional design changes the denominator across phases.

| Test                                 | Statistic                        | p-value     | Effect Size  | Interpretation                                                                                            |
|--------------------------------------|----------------------------------|-------------|--------------|-----------------------------------------------------------------------------------------------------------|
| Qwen augmentation (descriptive)      | 74.0% → 56.0% from L1 to L4      | –           | –            | Decline on the L0-conditional subset; not treated as confirmatory because selection depends on L0 success |
| Direction: CUDA↔OMP (McNemar)        | $n_{\text{paired}} = 24$ , exact | $p = 0.180$ | $ h  = 0.44$ | No significant asymmetry ( $\alpha = 0.01$ )                                                              |
| Direction: CUDA↔OpenCL (McNemar)     | $n_{\text{paired}} = 19$ , exact | $p = 1.000$ | $ h  = 0.46$ | No significant asymmetry ( $\alpha = 0.01$ )                                                              |
| Direction: OMP↔OpenCL (McNemar)      | $n_{\text{paired}} = 17$ , exact | $p = 0.289$ | $ h  = 0.53$ | No significant asymmetry ( $\alpha = 0.01$ )                                                              |
| Direction: CUDA↔OMP-target (McNemar) | $n_{\text{paired}} = 8$ , exact  | $p = 0.031$ | $ h  = 2.09$ | Large effect but not significant ( $\alpha = 0.01$ ; $n$ too small)                                       |
| Direction: OMP↔OMP-target (McNemar)  | $n_{\text{paired}} = 3$ , exact  | $p = 1.000$ | $ h  = 1.23$ | Underpowered ( $n = 3$ )                                                                                  |

## 1250 F Translation Prompt Template

1251 This appendix reproduces the translation prompt template used by the evaluation pipeline  
1252 (build\_translation\_prompt() in llm\_evaluate.py). All prompts follow this template; the  
1253 specific content—source code, build command, target API, and file lists—varies between tasks.

### 1254 F.1 System Message

1255 The system message follows the same template for all translation tasks, parameterized by the source  
1256 and target API names:

```
1257 You are a parallel programming expert
1258 specializing in {src_api} to {tgt_api}
1259 translation. Translate the provided source
1260 code to {tgt_api}. Output ONLY the translated
1261 code, no explanations. For each file, output
1262 a markdown code fence with the filename on
1263 the opening line:
1264
1265 ```{tgt_lang} filename={filename}
1266 <code>
1267 ```
1268
1269 Preserve the algorithm, I/O behavior, data
1270 formats, and output format exactly. The
1271 translated code must compile with the
1272 provided build command.
```

Listing 4: System message template. {src\_api} and {tgt\_api} are replaced with API display names (e.g., “CUDA”, “OpenMP”). {tgt\_lang} is the target language hint for the code fence (e.g., “cuda”, “cpp”, “c”).

### 1275 F.2 User Message Structure

1276 The user message is assembled from the following sections, in order:

- 1277 1. **Translation Task** — source and target API display names. The kernel name and benchmark  
1278 description are *not* included (anonymization).
- 1279 2. **Target Files to Produce** — list of genericized target filenames (e.g., translated\_0.cpp,  
1280 translated\_1.cl). When target infrastructure context is provided, an explanatory note  
1281 clarifies that only these files replace existing project files.
- 1282 3. **Build Command** — the anonymized compilation command in a code fence, so the LLM  
1283 can ensure API and flag compatibility.
- 1284 4. **Build Environment** — system dependencies from the target spec (e.g., “CUDA Toolkit  $\geq$   
1285 11.0”, “GCC  $\geq$  9.0”).
- 1286 5. **Source Code** — each source file presented as a numbered subsection (Source File 1,  
1287 Source File 2, ...) with all C/C++ comments stripped. The section heading includes the  
1288 source API name (e.g., “Source Code (CUDA)”).
- 1289 6. **Support / Header Files** (optional) — source-directory headers and code files, genericized  
1290 as Header File *N* and Code File *N*, with an instruction to inline definitions rather  
1291 than emit unresolvable #include directives. Code files are included up to a cumulative  
1292 50,000-character limit.
- 1293 7. **Target Infrastructure Context** (optional) — non-kernel target files (prompt payload entries  
1294 not in translation\_targets, plus target support headers) provided as read-only reference  
1295 so the LLM can match expected function signatures and data structures. Headed “DO NOT  
1296 MODIFY — for reference only” and genericized as Infrastructure File *N*. Omitted  
1297 when no non-kernel target files or target support headers remain after filtering.

1298 Listing 5 shows the complete user message template with placeholders matching the system message  
1299 convention (Listing 4).

```

1300
1301 ## Translation Task
1302 Source API: {src_api} -> Target API: {tgt_api}
1303
1304 ## Target Files to Produce
1305 - translated_0.{ext}
1306 - translated_1.{ext}
1307   [... one entry per translation target]
1308
1309 _These files will replace the corresponding
1310 files in the target project directory. All other
1311 project files (Makefile, headers, utilities)
1312 remain unchanged._
1313
1314 ## Build Command (your code must work with this)
1315 {anonymized_build_command}
1316
1317 ## Build Environment
1318 - {dependency_1}
1319 - {dependency_2}
1320
1321 ## Source Code ({src_api})
1322 ### Source File 1
1323 ‘‘{src_lang}
1324 {source_code_comments_stripped}
1325 Source File 2
1326
1327 {source_code_comments_stripped}
1328
1329 Support / Header Files (conditional)
1330
1331 These files exist in the source build directory
1332 but may NOT exist in the target directory. If
1333 your translated code needs definitions from them,
1334 inline the definitions directly rather than
1335 using #include.
1336
1337 Header File 1
1338
1339 {header_contents_comments_stripped}
1340
1341 Target Infrastructure Context (conditional)
1342
1343 (DO NOT MODIFY -- for reference only)
1344 These files exist in the target build directory
1345 and will NOT be modified. Your translated code
1346 must be compatible with them.
1347
1348 Infrastructure File 1
1349
1350 {infrastructure_code_comments_stripped}
1351

```

Listing 5: User message template. Sections marked (*conditional*) appear only when the translation task includes the relevant files. {src\_api} and {tgt\_api} are API display names; {src\_lang} and {tgt\_lang} are language hints for code fences. All source code has C/C++ comments stripped before inclusion.

### 1352 F.3 Anonymization Details

1353 Six anonymization measures are applied during prompt construction to reduce the risk that benchmark-  
1354 specific identifiers trigger memorized translations. These measures complement the AST-driven  
1355 source augmentation (Section 2.2): anonymization targets prompt *metadata*, while augmentation  
1356 modifies source code *structure*.

- 1357 1. **Kernel identity omission.** The kernel name and benchmark description are excluded from  
1358 the prompt entirely.

2. **Comment stripping.** All C/C++ line (`//`) and block (`/* */`) comments are removed from every file shown to the LLM—source files, support files, and target infrastructure files alike. The stripper uses a state-machine parser that preserves string, character, and raw string literals.
3. **Source and support filename genericization.** Source files are labeled Source File 1, Source File 2, etc. Support files are labeled Header File  $N$  or Code File  $N$  by extension. Target infrastructure files are labeled Infrastructure File  $N$ .
4. **Target filename genericization.** Target output filenames are replaced with `translated_0.ext`, `translated_1.ext`, etc., preserving original file extensions. An internal mapping restores original filenames when writing LLM output to disk, including any cross-file `#include` references the LLM emits using the generic names.
5. **Build command anonymization.** Kernel-specific identifiers are removed from the build command: `make target` is reduced to `make` (relying on the Makefile default target), and kernel names in other command strings are replaced with a generic placeholder.
6. **API-name retention.** The source and target API names (e.g., “CUDA”, “OpenMP”) are intentionally *not* anonymized, as they are essential to the translation task specification.

## G Evaluation Cost and Reproducibility

This appendix groups the practical information needed to reproduce the reported evaluation campaigns: model/provider settings, the shared hardware/software platform, and the aggregate token and wall-clock cost of the three campaigns. Table 18 records the provider-side model settings. Table 19 records the shared execution environment. Table 20 then reports aggregate campaign cost.

Table 18: Model configurations. <sup>†</sup>Azure reasoning models do not accept explicit temperature or top- $p$  parameters; sampling is controlled internally by the provider.

| Model              | Provider     | Arch.       | Parameters        | Reasoning                            | Temp.            |
|--------------------|--------------|-------------|-------------------|--------------------------------------|------------------|
| Qwen 3.5 397B-A17B | Together AI  | MoE         | 397B / 17B active | <code>enable_thinking</code>         | 0.7              |
| GPT-5.4            | Azure OpenAI | proprietary | proprietary       | <code>reasoning_effort=medium</code> | N/A <sup>†</sup> |
| GPT-5.3-codex      | Azure OpenAI | proprietary | proprietary       | <code>reasoning_effort=medium</code> | N/A <sup>†</sup> |

Table 19: Hardware and software configuration shared across all three model campaigns.

| Component      | Specification                                             |
|----------------|-----------------------------------------------------------|
| GPU            | NVIDIA GeForce RTX 4070 (12 GB, Ada Lovelace, sm_89)      |
| CPU            | AMD Ryzen 9 7900X (12 cores / 24 threads, 4.7 GHz)        |
| RAM            | 32 GB DDR5                                                |
| OS             | Ubuntu 24.04 LTS                                          |
| CUDA toolkit   | NVIDIA HPC SDK 24.3 (nvcc 12.3, V12.3.103)                |
| C/C++ compiler | GCC 12.4.0                                                |
| OpenMP         | GCC <code>-fopenmp</code> ; nvc++ 24.3 for target offload |
| OpenCL runtime | NVIDIA via HPC SDK 24.3                                   |
| Python         | 3.12.3                                                    |

The higher wall time for GPT-5.4 reflects its use of internal reasoning (`reasoning_effort=medium`), which approximately doubles per-call latency relative to the non-reasoning models at comparable output lengths.

Per-result token counts reflect the usage statistics returned by each provider’s API, which include system-prompt tokens. They may modestly undercount actual billed consumption due to provider-internal formatting tokens or SDK-level retries on transient failures. Researchers budgeting a replication should consult current provider pricing and allow a margin for such overhead. The Qwen 3.5 task count includes 82 records involving `KNOWN_FAIL` specs that are excluded from pass-rate denominators (Section B.4); the GPT-5.4 and GPT-5.3-codex batches pre-excluded such specs, so all on-disk records are evaluation-eligible.

Table 20: Evaluation campaign cost summary. All metrics are aggregated from per-result JSON metadata. Token counts reflect only the tokens recorded per evaluation call (see text).

| Metric                  | Qwen 3.5                                | GPT-5.4   | GPT-5.3-codex |
|-------------------------|-----------------------------------------|-----------|---------------|
| Tasks evaluated         | 708                                     | 822       | 814           |
| Total tokens            | 13.6M                                   | 14.9M     | 13.2M         |
| Input                   | 8.9M                                    | 9.3M      | 9.2M          |
| Output                  | 4.7M                                    | 5.6M      | 4.0M          |
| Avg. tokens per task    | 19,149                                  | 18,138    | 16,204        |
| Total LLM wall time (h) | 14.0                                    | 27.4      | 11.9          |
| Campaign dates (2026)   | Apr 21–24                               | Apr 25–27 | Apr 30–May 1  |
| <b>All models</b>       | 2,344 tasks, 41.7M tokens, 53.4 h total |           |               |

## H Evaluation Card

Following Model Cards Mitchell et al. [2019] and Datasheets Gebru et al. [2021], we provide an *Evaluation Card* for PARBENCH that states what it measures, what it does not measure, the claims it supports, the assumptions it relies on, and how results should be reported. Table 21 summarizes the scope; the subsections below justify each point.

Table 21: Evaluation Card summary. Each row states a key property of PARBENCH’s measurement scope. Detailed discussion appears in the referenced subsection.

| Dimension                             | Key Points                                                                                          | Ref. |
|---------------------------------------|-----------------------------------------------------------------------------------------------------|------|
| <i>What PARBENCH measures</i>         |                                                                                                     |      |
| Declared-oracle pass/-fail            | Binary pass/fail via conjunctive build→run→verify pipeline                                          | §H.1 |
| Failure taxonomy                      | Four failure classifications (EXTRACT_FAIL, BUILD_FAIL, RUN_FAIL, VERIFY_FAIL) plus PASS            | §H.1 |
| Direction asymmetry                   | Pass-rate variation across 10 translation directions (3 APIs + OMP target)                          | §H.1 |
| Surface-form robustness               | Stability under six AST-level source transforms at four levels                                      | §H.1 |
| <i>What PARBENCH does not measure</i> |                                                                                                     |      |
| Performance                           | No speedup, throughput, occupancy, memory, or energy measurement                                    | §H.2 |
| Numerical accuracy                    | 80 of 87 non-KNOWN_FAIL specs use weak oracles (exit code + stdout pattern)                         | §H.2 |
| Code quality                          | No readability, maintainability, or style assessment of translations                                | §H.2 |
| Repo-level translation                | Kernel-centric by design; host code and build infra remain fixed                                    | §H.2 |
| Self-repair capability                | Single-attempt evaluation; iterative feedback not exercised                                         | §H.2 |
| <i>Validity and assumptions</i>       |                                                                                                     |      |
| Valid claims                          | Pass@k, failure taxonomy, direction asymmetry, surface-form robustness on the L0-conditional subset | §H.3 |
| Invalid claims                        | “Efficient code,” “production-ready,” “understands parallelism,” “not memo-rizing”                  | §H.4 |
| Key assumptions                       | Baseline-valid source specs, oracle sufficiency (weakest), single-platform generalization           | §H.5 |
| <i>Guidance</i>                       |                                                                                                     |      |
| User guidance                         | Adding models/kernels, interpreting results, comparing models, common pitfalls                      | §H.6 |
| Reporting protocol                    | Required and recommended items for publications using PARBENCH                                      | §H.7 |

### H.1 What PARBENCH Measures

PARBENCH measures the following properties of LLM-based parallel code translation:

- 1. Binary declared-oracle pass/fail of kernel-level translation.** The harness evaluates whether LLM-translated code compiles, executes, and produces output matching the spec’s verification strategies. Verification applies a conjunction of strategies: all declared checks

- (exit code, stdout pattern, and optionally numeric comparison or file hash) must pass for a record to receive PASS status. The result is PASS or one of four failure statuses (EXTRACT\_FAIL, BUILD\_FAIL, RUN\_FAIL, VERIFY\_FAIL); timeouts are classified under the relevant pipeline stage (BUILD\_FAIL or RUN\_FAIL).
2. **Translation capability across three primary APIs and ten directions.** CUDA, OpenMP, and OpenCL form six bidirectional standard directions; OpenMP Target adds four case-study directions. The evaluation matrix covers 142 unique source–target pairs (Section 2.3).
  3. **Failure mode taxonomy.** Four failure classifications (EXTRACT\_FAIL, BUILD\_FAIL, RUN\_FAIL, VERIFY\_FAIL) enable diagnostic analysis of *where* in the pipeline translations fail. Build failures indicate incomplete API-surface adaptation; run failures indicate runtime errors or timeouts; verify failures indicate incorrect output from otherwise executable code. This granularity distinguishes PARBENCH from compile-only or text-similarity evaluations.
  4. **Surface-form robustness on the L0-conditional subset.** Six AST-level transforms (Swap-Condition, ArithmeticTransform, ChangeNames, TypedefExpansion, PointerArithmetic-ToArrayIndex, ChangeFunctionNames) at four intensity levels (L1–L4) test whether pass/fail status is stable under AST-level code changes intended to preserve behavior (Section 2.2). The augmentation campaign applies only to direction–kernel pairs that qualify under the L0-conditional filter.
  5. **Direction asymmetry.** Pass-rate variation across translation directions quantifies structural difficulty differences between API pairs. CUDA-to-OpenMP consistently passes at higher rates than OpenCL-to-CUDA across all three evaluated models (Section 5.3).
  6. **Cross-model discrimination.** The benchmark distinguishes models that differ in translation capability, subject to sampling-condition caveats (Section H.5). In the current evaluation, pairwise analysis on 142 shared tasks yields Cohen’s  $h \approx 0.8$  between Qwen 3.5 and both GPT models ( $|h| = 0.80$  for both GPT-5.4 and GPT-5.3-codex, reflecting their identical overall pass rates), while GPT-5.4 and GPT-5.3-codex are statistically indistinguishable ( $h = 0.00$ , McNemar  $p = 0.72$ ).
  7. **Per-kernel difficulty heterogeneity.** For Qwen 3.5, canonical pass rates range from 72.2% (floydwarshall, iso2dfd) to 0% (10 kernels), enabling fine-grained analysis of which computational patterns—stencils, graph traversal, ODE integration, pointer-heavy data structures—are harder to translate (Appendix Table 24 reports rates including augmentation levels).

## 1431 H.2 What PARBENCH Does Not Measure

1432 The following properties are explicitly outside PARBENCH’s measurement scope:

1. **Performance and resource efficiency.** No speedup, throughput, occupancy, bandwidth, memory, or energy measurement is performed. Wall-clock time is captured but is unreliable for sub-millisecond baselines. No kernel-level profiling (ncu/nsys for CUDA, omp\_get\_wtime for OpenMP) is included. A harness-passing translation that runs 100× slower than the original receives PASS. TRACY Gong et al. [2025] addresses this gap.
2. **Code quality, readability, or maintainability.** Translated code is stored in result files but no AST analysis, cyclomatic complexity, or style metrics are computed. A harness-passing but unreadable translation receives PASS.
3. **Partial pass credit or oracle-strength gradations.** Verification is binary: a numerical result that deviates from the reference by 0.02% beyond the configured tolerance receives VERIFY\_FAIL, identical to a completely wrong output. There is no partial-credit classification.
4. **Numerical accuracy for most specs.** Only 7 of 87 non-KNOWN\_FAIL specs (8%) have medium or strong oracles (numeric\_comparison or file\_hash). The remaining 80 specs (92%) rely on exit\_code and stdout\_pattern, which verify that the program runs and prints expected banners but do not independently validate numerical results. A translation that computes wrong numbers but prints the expected completion message would receive PASS on those specs. The oracle strength distribution across all 206 specs is: 2 strong (file\_hash), 5 medium (numeric\_comparison), 46 weak (stdout\_pattern + exit\_code), and 153 untagged (effectively weak). Upgrading weak oracles to numeric

comparison requires per-kernel reference-output engineering; we prioritize this in future work.

5. **Iterative self-repair or agentic translation.** Each sample is a single LLM call with no feedback loop. The evaluation pipeline implements iterative repair (multi-turn error feedback with linker analysis via `--max-retries`), but all reported results use `max_retries=1` (zero-shot). Evaluating multi-turn repair strategies—where build or verification failures are fed back to the model for corrected attempts—and agentic translation pipelines that combine planning, tool use, and iterative refinement remain future work.
6. **Cross-platform portability.** All evaluations run on a single platform: NVIDIA RTX 4070 (sm\_89), AMD Ryzen 9 7900X, Ubuntu 24.04, HPC SDK 24.3 (Appendix Table 19). No testing on other GPU architectures (A100, H100), CPU-only environments, or other OS configurations.
7. **Repository-level translation.** By design, PARBENCH isolates kernel translation from build-system reconstruction. Host code, Makefiles, and I/O routines remain fixed. The benchmark does not measure the LLM’s ability to restructure project files, generate build systems, or coordinate multi-component applications (cf. ParEval-Repo Davis et al. [2025]).
8. **Training-data memorization beyond surface form (definitively).** Augmentation tests surface-form robustness, not algorithmic memorization (Section H.4).

### 1471 H.3 Valid Claims

1472 The following conclusions can be drawn from PARBENCH results when accompanied by the stated  
1473 conditions:

- 1474 1. **“Model X achieves Y% pass@ $k$  on kernel-centric parallel code translation across  $Z$   
1475 directions.”** Valid when: the exact pass@ $k$  metric is specified (pass@1 vs. pass@3), aug-  
1476 mentation levels are stated (L0 only vs. L0–L4), directions are enumerated, and confidence  
1477 intervals are included.
- 1478 2. **“Build-stage failure is the dominant failure mode for Model X.”** Valid when: the failure  
1479 taxonomy breakdown is reported with counts and percentages from the full record set.
- 1480 3. **“Direction A→B is harder/easier than Direction C→D for Model X.”** Valid when: backed  
1481 by per-direction pass rates with Wilson confidence intervals. Claims of statistical signifi-  
1482 cance should use McNemar’s test on paired kernels with appropriate multiple-comparison  
1483 correction.
- 1484 4. **“Model X maintains declared-oracle pass rates across L1–L4 on the L0-conditional  
1485 subset.”** Valid when: the L0-conditional filter is disclosed, the qualifying subset size is  
1486 stated, the direction restriction is noted, and the analysis is presented as descriptive rather  
1487 than as a definitive memorization test.
- 1488 5. **“Model X discriminates from Model Y on PARBENCH.”** Valid when: (a) all compared  
1489 models are evaluated on the same task set, (b) McNemar’s paired test is reported with effect  
1490 size and concordance table, (c) sampling-configuration differences are explicitly disclosed if  
1491 temperatures or reasoning modes differ, and (d) the claim is scoped to “PARBENCH tasks,”  
1492 not “parallel translation in general.”
- 1493 6. **“Kernel K is harder to translate than Kernel J.”** Valid when: both kernels have sufficient  
1494 sample sizes ( $n \geq 18$  recommended) and confidence intervals are reported.

### 1495 H.4 Invalid Claims

1496 The following conclusions **cannot** be drawn from PARBENCH results:

- 1497 1. **“Model X produces efficient/fast parallel code.”** Performance is not measured. A PASS  
1498 means the code compiles, runs, and satisfies the declared oracle—not that it runs at accept-  
1499 able speed.
- 1500 2. **“Model X’s translations are production-ready.”** No code review, security audit, or race-  
1501 condition detection is performed. Parallel code can have latent data races that produce  
1502 correct output on some executions.

- 1503 3. **“Model X understands parallel programming.”** PARBENCH measures behavioral out-  
1504 comes, not internal representations. A model could pass by pattern matching without  
1505 “understanding” synchronization semantics.
- 1506 4. **“PARBENCH pass rate generalizes to all HPC translation tasks.”** The corpus contains  
1507 87 non-KNOWN\_FAIL specs from five suites, with Rodinia contributing 53 (61%). Domains  
1508 such as distributed-memory MPI, GPU tensor operations, and FPGA HLS are outside its  
1509 scope.
- 1510 5. **“Model X is definitively better than Model Y at parallel translation.”** Only valid under  
1511 matched sampling conditions. If temperatures, reasoning modes, or provider-side controls  
1512 differ (as they do between Qwen 3.5 and the Azure models in this paper), the observed gap  
1513 reflects an unknown mixture of model capability and sampling configuration (Section B.4).
- 1514 6. **“Surface-form robustness proves the model is not memorizing.”** Surface-form aug-  
1515 mentation (variable renaming, syntax sugar) tests robustness to cosmetic code changes.  
1516 Algorithmic memorization—where the model recognizes the computation and produces a  
1517 previously seen translation template—is not tested by these transforms.
- 1518 7. **“A PASS means the translation is numerically faithful.”** For the 80 of 87 specs with  
1519 weak oracles, verification checks that the program ran and printed expected output banners.  
1520 Numerical results are not independently validated. Only 7 specs have medium or strong  
1521 oracles that verify numerical output.
- 1522 8. **“OpenCL→CUDA is impossible for LLMs.”** Qwen 3.5 achieves 0% but both GPT-5.4 and  
1523 GPT-5.3-codex achieve 19.3% on the same direction. Zero-rate results for a single model  
1524 reflect that model’s capability, not an inherent impossibility.

## 1525 H.5 Assumptions

1526 PARBENCH’s evaluation results are meaningful under the following assumptions. We note mitigations  
1527 and residual risks for each.

- 1528 1. **Source implementations are baseline-valid on the reference platform.** PARBENCH treats  
1529 the original benchmark implementations as the reference baseline. If a source implemen-  
1530 tation has a latent bug, faithful translations of that bug would pass the declared oracle.  
1531 *Mitigation:* all 87 non-KNOWN\_FAIL specs pass the baseline build/run/verify pipeline on  
1532 the reference platform.
- 1533 2. **Declared oracles are sufficient to detect material translations.** This is the weakest  
1534 assumption. For the 80 specs with weak oracles, a translation that produces wrong numerical  
1535 results but expected program flow (runs to completion, prints expected banners) would  
1536 receive a false PASS. *Mitigation:* 7 specs have numeric\_comparison or file\_hash  
1537 oracles; the dominant failure mode (BUILD\_FAIL) is caught by all oracle types regardless  
1538 of oracle strength. Upgrading weak oracles is identified as future work (Section 6).
- 1539 3. **The single evaluation platform generalizes to other NVIDIA GPU configurations.** All  
1540 results are from one machine (RTX 4070, sm\_89). Different GPU architectures may have  
1541 different compiler behavior, runtime characteristics, or memory limits.
- 1542 4. **Three samples per task capture meaningful stochastic variance under each provider’s**  
1543 **sampling regime.** For Qwen 3.5 this uses temperature 0.7; for GPT-5.4 and GPT-5.3-codex  
1544 sampling is provider-controlled (Section B.4). The pass@1-to-pass@3 gap (11.3 pp for  
1545 Qwen 3.5, 7.0 pp for GPT-5.4, 5.6 pp for GPT-5.3-codex) suggests moderate within-task  
1546 variance. More samples would provide tighter estimates at proportional cost.
- 1547 5. **Kernel-centric translation isolates parallel programming capability.** By fixing host code  
1548 and build infrastructure, PARBENCH prevents build-system reconstruction from dominating  
1549 results (cf. ParEval-Repo’s 0% at >133 SLoC Davis et al. [2025]). However, this also means  
1550 that a model’s ability to handle multi-component integration—a real-world requirement—is  
1551 explicitly not tested.
- 1552 6. **The 9 KNOWN\_FAIL exclusions are legitimately infrastructure failures.** Each exclusion  
1553 is documented with a specific technical cause (CUDA 12 API deprecation, missing libraries,  
1554 pre-existing build or runtime errors).

- 1555 7. **Benchmark suites are representative of real HPC workloads.** The five suites are well-  
 1556 established in HPC research (Rodinia: IISWC 2009 Che et al. [2009], HeCBench: 2023 Jin  
 1557 and Vetter [2023], XSBench: ANL Tramm et al. [2014]). However, they over-represent  
 1558 structured stencil and reduction patterns and under-represent irregular computations, multi-  
 1559 physics coupling, and modern GPU programming idioms (tensor cores, cooperative groups).
- 1560 8. **Augmentation transforms are intended to preserve source behavior.** Each transform is  
 1561 backed by libclang AST analysis and validated by 15 unit tests. 80 of 87 specs pass all  
 1562 L1–L4 augmented baseline verification. The 7 failures are `omp_target`-specific: condition  
 1563 operand swapping (e.g.,  $a < b \rightarrow b > a$ ) causes the GPU-offloading compiler to generate  
 1564 device code that crashes or produces divergent output, even though the transform is behavior-  
 1565 preserving under standard C compilation. Variable renaming is disabled for files containing  
 1566 OpenMP pragmas; the 5 HeCBench failures involve only condition swapping, while the  
 1567 2 remaining failures (RSBench, XSBench at L4) co-occur with additional transforms on  
 1568 non-pragma helper files. These 7 augmented variants fail their own baseline verification,  
 1569 meaning the behavior-preserving assumption cannot be confirmed at those levels; they are  
 1570 flagged in analysis but still participate as source text in LLM evaluation (since the source is  
 1571 used as prompt input, not compiled).
- 1572 9. **Prompt anonymization does not change task difficulty.** Stripping comments, generi-  
 1573 cizing filenames, and removing kernel names reduce memorization cues but might also  
 1574 remove helpful context, potentially making PARBENCH results conservative relative to  
 1575 non-anonymized use.

## 1576 H.6 User Guidance

1577 **Adding a new model.** Register the model in the evaluation pipeline’s model registry, specifying  
 1578 provider, API model identifier, and thinking/temperature configuration. Run the canonical campaign  
 1579 with all suites; the evaluation pipeline handles prompt construction, LLM calls, and file extraction,  
 1580 and the harness handles build/run/verify.

1581 **Adding new kernels.** Write a spec JSON following the schema documented in Appendix B.1. The  
 1582 spec must define identity (unique ID, API, suite), file partitioning (prompt payload, translation targets,  
 1583 support files), build commands, run configuration, and verification strategies. Validate the baseline by  
 1584 running the harness verify command. Append to the manifest (append-only).

1585 **Interpreting pass rates.** Always report the denominator (number of evaluation-eligible records  
 1586 after KNOWN\_FAIL exclusion), the augmentation level (L0 vs. L0–L4), and Wilson 95% confi-  
 1587 dence intervals. Per-direction breakdowns are essential—aggregate rates conceal strong direction  
 1588 asymmetry.

1589 **Comparing models.** Use paired tests on task outcomes (same tasks, same direction–kernel pairs):  
 1590 Fisher’s exact test for pairwise comparisons, McNemar’s test for concordance analysis. Report  
 1591 concordance tables (both-pass, both-fail, discordant). Always disclose temperature, reasoning mode,  
 1592 and any provider-side sampling controls. If sampling conditions are unmatched, state this explicitly  
 1593 and scope the comparison accordingly.

1594 **Augmentation analysis.** Use the L0-conditional filter: include a pair in the augmentation campaign  
 1595 only if  $\geq 1$  of its L0 samples passes. Report the filter rate (fraction of pairs qualifying). Results  
 1596 apply only to the qualifying subset. Single samples per augmentation level (L1–L4) trade statistical  
 1597 precision for budget efficiency.

### 1598 Common pitfalls.

- 1599 • Do not compare aggregate pass rates across models without disclosing sampling conditions.
- 1600 • Do not claim “the model understands parallel programming” from pass rates alone.
- 1601 • Do not interpret weak-oracle PASS as evidence of numerical fidelity.
- 1602 • Do not modify the KNOWN\_FAIL exclusion policy without documenting the change and its
- 1603 effect on denominators.
- 1604 • Do not change spec run arguments without reading the source code’s `argc` check—a
- 1605 documented cause of past silent failures.

## 1606 H.7 Recommended Reporting Protocol

1607 When publishing results obtained with PARBENCH, we recommend reporting at minimum the items  
 1608 in Table 22. For cross-model comparisons and augmentation analyses, the additional items in Table 23  
 1609 strengthen interpretability.

Table 22: Required reporting items for PARBENCH results.

| Item                 | What to report                                                                                            |
|----------------------|-----------------------------------------------------------------------------------------------------------|
| Model identity       | Provider, model ID, version or date                                                                       |
| Sampling config      | Temperature/top- $p$ if caller-controlled, otherwise “provider-controlled”; reasoning mode; $n$ (samples) |
| pass@ $k$            | pass@1 and pass@3 with confidence intervals                                                               |
| Failure taxonomy     | Counts and percentages for each failure type                                                              |
| Direction break-down | Per-direction pass rates with Wilson CIs                                                                  |
| Denominator          | Total evaluation-eligible records after KNOWN_FAIL exclusion                                              |
| Augmentation scope   | Levels evaluated, filter criterion, subset size                                                           |
| Platform             | GPU, CPU, OS, compiler versions                                                                           |

Table 23: Recommended additional reporting items.

| Item                           | What to report                                                |
|--------------------------------|---------------------------------------------------------------|
| <i>Cross-model comparisons</i> |                                                               |
| Paired test                    | McNemar on paired tasks with concordance table                |
| Effect size                    | Cohen’s $h$ for overall and per-direction gaps                |
| Sampling caveat                | Explicit statement if conditions are unmatched                |
| Task variance                  | Fraction of hard-fail ( $0/n$ ), noisy, always-pass ( $n/n$ ) |
| <i>Augmentation analysis</i>   |                                                               |
| Per-level rates                | Descriptive rates at L0–L4 on balanced subset                 |
| Filter disclosure              | “X of 142 pairs qualify (Y%)”                                 |
| Direction scope                | Which direction(s) the analysis covers                        |
| Power analysis                 | Detectable effect size at 80% power                           |

## 1610 I Artifact Availability

1611 An anonymized artifact snapshot is submitted with this paper and contains the curated PARBENCH  
 1612 specifications, the build-run-verify harness, the evaluation and augmentation pipelines, and per-record  
 1613 result JSONs for all three evaluated models. The submitted artifact contains 206 generated JSON  
 1614 specifications, including the 96 curated specifications used in the paper and the 87 eval-eligible  
 1615 specifications that pass the baseline build-run-verify pipeline. Each kernel specification embeds  
 1616 its provenance: the upstream suite, repository URL, pinned commit hash, license, build recipe,  
 1617 run configuration, and verification oracle (Section 3.1). If accepted, the same artifact will be de-  
 1618 anonymized, mirrored to GitHub under the MIT license for the PARBENCH framework, and archived  
 1619 on Zenodo to provide a persistent DOI; upstream suites retain their original licenses: Rodinia (BSD-  
 1620 like), XSBench (MIT), RSBench (MIT), HeCBench (BSD-3-Clause), and mixbench (GPL-2.0).  
 1621 Structured metadata is provided now through the JSON specification schema and per-spec provenance  
 1622 fields. Because PARBENCH is an executable evaluation benchmark rather than a model-training  
 1623 dataset, Croissant metadata is not applicable here.

## 1624 J Supplementary Tables and Figures

1625 This appendix retains the dense reference figures and full per-kernel tables that are easier to consult  
 1626 separately from the narrative appendices above. Per-kernel translation heatmaps appear in Figure 2  
 1627 (main body). The figures below focus on direction-level pass@ $k$  behavior, cross-suite comparisons,  
 1628 and full per-kernel rate tables.




Figure 18: Pass@ $k$  by translation direction (Qwen 3.5, L0, three samples per task).




Figure 19: Pass@ $k$  by translation direction (GPT-5.4, L0, three samples per task).




Figure 20: Pass@ $k$  by translation direction (GPT-5.3-codex, L0, three samples per task).




Figure 21: Cross-suite pass rate comparison (L0, Qwen 3.5). Per-suite aggregate pass rates with Wilson 95% CIs across all five benchmark suites.




Figure 22: Cross-suite pass rate comparison (L0, GPT-5.4). Per-suite aggregate pass rates with Wilson 95% CIs across all five benchmark suites.




Figure 23: Cross-suite pass rate comparison (L0, GPT-5.3-codex). Per-suite aggregate pass rates with Wilson 95% CIs across all five benchmark suites.

Table 24: Full per-kernel pass rates across all 31 evaluated kernels (Qwen 3.5, 626 evaluation-eligible records including L0 and augmentation levels, 31 kernels across 5 suites). Results involving KNOWN\_FAIL specs excluded.

| Suite    | Kernel         | Total | PASS | Fail | Rate  | 95% Wilson CI  |
|----------|----------------|-------|------|------|-------|----------------|
| HeCBench | convolution1d  | 10    | 2    | 8    | 20.0% | [5.7%, 51.0%]  |
| HeCBench | floydwarshall  | 38    | 29   | 9    | 76.3% | [60.8%, 87.0%] |
| HeCBench | heat2d         | 38    | 29   | 9    | 76.3% | [60.8%, 87.0%] |
| HeCBench | iso2dfd        | 38    | 32   | 6    | 84.2% | [69.6%, 92.6%] |
| HeCBench | jacobi         | 10    | 3    | 7    | 30.0% | [10.8%, 60.3%] |
| HeCBench | md             | 10    | 4    | 6    | 40.0% | [16.8%, 68.7%] |
| HeCBench | nqueen         | 10    | 6    | 4    | 60.0% | [31.3%, 83.2%] |
| HeCBench | page-rank      | 10    | 6    | 4    | 60.0% | [31.3%, 83.2%] |
| HeCBench | scan           | 6     | 0    | 6    | 0.0%  | [0.0%, 39.0%]  |
| HeCBench | stencil1d      | 14    | 9    | 5    | 64.3% | [38.8%, 83.7%] |
| Rodinia  | backprop       | 6     | 0    | 6    | 0.0%  | [0.0%, 39.0%]  |
| Rodinia  | bfs            | 34    | 15   | 19   | 44.1% | [28.9%, 60.5%] |
| Rodinia  | bptree         | 18    | 0    | 18   | 0.0%  | [0.0%, 17.6%]  |
| Rodinia  | cf             | 22    | 6    | 16   | 27.3% | [13.2%, 48.2%] |
| Rodinia  | dwt2d          | 6     | 0    | 6    | 0.0%  | [0.0%, 39.0%]  |
| Rodinia  | gaussian       | 6     | 0    | 6    | 0.0%  | [0.0%, 39.0%]  |
| Rodinia  | heartwall      | 18    | 0    | 18   | 0.0%  | [0.0%, 17.6%]  |
| Rodinia  | hotspot        | 30    | 14   | 16   | 46.7% | [30.2%, 63.9%] |
| Rodinia  | hotspot3d      | 34    | 20   | 14   | 58.8% | [42.2%, 73.6%] |
| Rodinia  | lavamd         | 18    | 0    | 18   | 0.0%  | [0.0%, 17.6%]  |
| Rodinia  | lud            | 30    | 14   | 16   | 46.7% | [30.2%, 63.9%] |
| Rodinia  | myocyte        | 18    | 0    | 18   | 0.0%  | [0.0%, 17.6%]  |
| Rodinia  | nn             | 6     | 0    | 6    | 0.0%  | [0.0%, 39.0%]  |
| Rodinia  | nw             | 30    | 11   | 19   | 36.7% | [21.9%, 54.5%] |
| Rodinia  | particlefilter | 26    | 9    | 17   | 34.6% | [19.4%, 53.8%] |
| Rodinia  | pathfinder     | 30    | 8    | 22   | 26.7% | [14.2%, 44.4%] |
| Rodinia  | sradi          | 26    | 9    | 17   | 34.6% | [19.4%, 53.8%] |
| Rodinia  | streamcluster  | 22    | 1    | 21   | 4.5%  | [0.8%, 21.8%]  |
| mixbench | mixbench       | 22    | 2    | 20   | 9.1%  | [2.5%, 27.8%]  |
| RSBench  | rsbench        | 18    | 0    | 18   | 0.0%  | [0.0%, 17.6%]  |
| XSBench  | xsbench        | 22    | 1    | 21   | 4.5%  | [0.8%, 21.8%]  |

Table 25: Full per-kernel pass rates across all 31 evaluated kernels (GPT-5.4, 822 evaluation-eligible records including L0 and augmentation levels, 31 kernels across 5 suites). KNOWN\_FAIL specs pre-excluded from evaluation batch.

| Suite    | Kernel         | Total | PASS | Fail | Rate   | 95% Wilson CI   |
|----------|----------------|-------|------|------|--------|-----------------|
| HeCBench | convolution1d  | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | floydwarshall  | 42    | 41   | 1    | 97.6%  | [87.7%, 99.6%]  |
| HeCBench | heat2d         | 42    | 40   | 2    | 95.2%  | [84.2%, 98.7%]  |
| HeCBench | iso2dfd        | 42    | 42   | 0    | 100.0% | [91.6%, 100.0%] |
| HeCBench | jacobi         | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | md             | 14    | 13   | 1    | 92.9%  | [68.5%, 98.7%]  |
| HeCBench | nqueen         | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | page-rank      | 14    | 12   | 2    | 85.7%  | [60.1%, 96.0%]  |
| HeCBench | scan           | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | stencil1d      | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| Rodinia  | backprop       | 10    | 7    | 3    | 70.0%  | [39.7%, 89.2%]  |
| Rodinia  | bfs            | 38    | 35   | 3    | 92.1%  | [79.2%, 97.3%]  |
| Rodinia  | bptree         | 26    | 13   | 13   | 50.0%  | [32.1%, 67.9%]  |
| Rodinia  | cfid           | 34    | 28   | 6    | 82.4%  | [66.5%, 91.7%]  |
| Rodinia  | dwt2d          | 10    | 7    | 3    | 70.0%  | [39.7%, 89.2%]  |
| Rodinia  | gaussian       | 6     | 0    | 6    | 0.0%   | [0.0%, 39.0%]   |
| Rodinia  | heartwall      | 22    | 6    | 16   | 27.3%  | [13.2%, 48.2%]  |
| Rodinia  | hotspot        | 34    | 28   | 6    | 82.4%  | [66.5%, 91.7%]  |
| Rodinia  | hotspot3d      | 38    | 26   | 12   | 68.4%  | [52.5%, 80.9%]  |
| Rodinia  | lavamd         | 26    | 8    | 18   | 30.8%  | [16.5%, 50.0%]  |
| Rodinia  | lud            | 34    | 27   | 7    | 79.4%  | [63.2%, 89.7%]  |
| Rodinia  | myocyte        | 22    | 5    | 17   | 22.7%  | [10.1%, 43.4%]  |
| Rodinia  | nn             | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| Rodinia  | nw             | 34    | 23   | 11   | 67.6%  | [50.8%, 80.9%]  |
| Rodinia  | particlefilter | 42    | 33   | 9    | 78.6%  | [64.1%, 88.3%]  |
| Rodinia  | pathfinder     | 34    | 27   | 7    | 79.4%  | [63.2%, 89.7%]  |
| Rodinia  | srad           | 38    | 28   | 10   | 73.7%  | [58.0%, 85.0%]  |
| Rodinia  | streamcluster  | 22    | 7    | 15   | 31.8%  | [16.4%, 52.7%]  |
| mixbench | mixbench       | 38    | 27   | 11   | 71.1%  | [55.2%, 83.0%]  |
| RSBench  | rsbench        | 38    | 29   | 9    | 76.3%  | [60.8%, 87.0%]  |
| XSBench  | xsbench        | 38    | 25   | 13   | 65.8%  | [49.9%, 78.8%]  |

Table 26: Full per-kernel pass rates across all 31 evaluated kernels (GPT-5.3-codex, 814 evaluation-eligible records including L0 and augmentation levels, 31 kernels across 5 suites). KNOWN\_FAIL specs pre-excluded from evaluation batch.

| Suite    | Kernel         | Total | PASS | Fail | Rate   | 95% Wilson CI   |
|----------|----------------|-------|------|------|--------|-----------------|
| HeCBench | convolution1d  | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | floydwarshall  | 42    | 40   | 2    | 95.2%  | [84.2%, 98.7%]  |
| HeCBench | heat2d         | 42    | 39   | 3    | 92.9%  | [81.0%, 97.5%]  |
| HeCBench | iso2dfd        | 42    | 42   | 0    | 100.0% | [91.6%, 100.0%] |
| HeCBench | jacobi         | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | md             | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | nqueen         | 14    | 13   | 1    | 92.9%  | [68.5%, 98.7%]  |
| HeCBench | page-rank      | 14    | 13   | 1    | 92.9%  | [68.5%, 98.7%]  |
| HeCBench | scan           | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| HeCBench | stencil1d      | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| Rodinia  | backprop       | 6     | 0    | 6    | 0.0%   | [0.0%, 39.0%]   |
| Rodinia  | bfs            | 38    | 35   | 3    | 92.1%  | [79.2%, 97.3%]  |
| Rodinia  | bptree         | 30    | 13   | 17   | 43.3%  | [27.4%, 60.8%]  |
| Rodinia  | cfid           | 34    | 27   | 7    | 79.4%  | [63.2%, 89.7%]  |
| Rodinia  | dwt2d          | 10    | 6    | 4    | 60.0%  | [31.3%, 83.2%]  |
| Rodinia  | gaussian       | 6     | 0    | 6    | 0.0%   | [0.0%, 39.0%]   |
| Rodinia  | heartwall      | 22    | 3    | 19   | 13.6%  | [4.7%, 33.3%]   |
| Rodinia  | hotspot        | 34    | 28   | 6    | 82.4%  | [66.5%, 91.7%]  |
| Rodinia  | hotspot3d      | 38    | 33   | 5    | 86.8%  | [72.7%, 94.2%]  |
| Rodinia  | lavamd         | 22    | 6    | 16   | 27.3%  | [13.2%, 48.2%]  |
| Rodinia  | lud            | 34    | 28   | 6    | 82.4%  | [66.5%, 91.7%]  |
| Rodinia  | myocyte        | 22    | 7    | 15   | 31.8%  | [16.4%, 52.7%]  |
| Rodinia  | nn             | 14    | 14   | 0    | 100.0% | [78.5%, 100.0%] |
| Rodinia  | nw             | 38    | 26   | 12   | 68.4%  | [52.5%, 80.9%]  |
| Rodinia  | particlefilter | 38    | 29   | 9    | 76.3%  | [60.8%, 87.0%]  |
| Rodinia  | pathfinder     | 34    | 28   | 6    | 82.4%  | [66.5%, 91.7%]  |
| Rodinia  | srad           | 34    | 26   | 8    | 76.5%  | [60.0%, 87.6%]  |
| Rodinia  | streamcluster  | 22    | 7    | 15   | 31.8%  | [16.4%, 52.7%]  |
| mixbench | mixbench       | 42    | 32   | 10   | 76.2%  | [61.5%, 86.5%]  |
| RSBench  | rsbench        | 38    | 20   | 18   | 52.6%  | [37.3%, 67.5%]  |
| XSBench  | xsbench        | 34    | 19   | 15   | 55.9%  | [39.5%, 71.1%]  |

## 1629 NeurIPS Paper Checklist

### 1630 1. Claims

1631 Question: Do the main claims made in the abstract and introduction accurately reflect the  
1632 paper’s contributions and scope?

1633 Answer: [Yes]

1634 Justification: The abstract and introduction define the benchmark contribution and  
1635 scope, while Section 5 reports L0 pass@1/pass@3 results (Qwen 23.9%/35.2%, GPT-  
1636 5.4 62.7%/69.7%, GPT-5.3-codex 62.7%/68.3%), failure taxonomy, direction asymmetry,  
1637 and augmentation findings. All numerical claims are verified against raw result JSONs.  
1638 Claims are scoped to the three evaluated models, declared-oracle PASS, and the tested  
1639 L0/L0-conditional protocols.

### 1640 2. Limitations

1641 Question: Does the paper discuss the limitations of the work performed by the authors?

1642 Answer: [Yes]

1643 Justification: Section 6 includes a dedicated limitations subsection addressing: three-model  
1644 scope, corpus coverage, single-GPU evaluation platform, weak oracle coverage, unmatched  
1645 sampling conditions across providers, and the absence of performance (runtime) evaluation.

### 1646 3. Theory

1647 Question: For each theoretical result, does the paper provide the full set of assumptions and  
1648 a complete proof?

1649 Answer: [N/A]

1650 Justification: This is an empirical benchmark paper with no theoretical results.

### 1651 4. Experiments

1652 Question: Does the paper fully disclose all the information needed to reproduce the main ex-  
1653 perimental results of the paper to the extent that it affects the main claims and/or conclusions  
1654 of the paper?

1655 Answer: [Yes]

1656 Justification: Section 4, Appendix G, and Appendix B provide model configurations (Ta-  
1657 ble 18), hardware specifications (Table 19), the complete prompt template (Appendix F),  
1658 and augmentation level definitions (Table 8). The evaluation pipeline is deterministic given  
1659 a fixed spec and model response.

### 1660 5. Code and data

1661 Question: Does the paper provide instructions for reproducing the main experimental results,  
1662 including code and data?

1663 Answer: [Yes]

1664 Justification: PARBENCH is included as anonymized supplementary material containing all  
1665 206 kernel specifications, the build-run-verify harness, the evaluation pipeline, augmentation  
1666 engine, and per-record result JSONs for all three models.

### 1667 6. Experimental details

1668 Question: Does the paper specify all the training and test details necessary to understand the  
1669 results?

1670 Answer: [Yes]

1671 Justification: No training is performed. Evaluation details are specified: temperature (0.7  
1672 for Qwen 3.5; provider-controlled for GPT-5.4 and GPT-5.3-codex), number of samples  
1673 (3), augmentation levels (L0–L4), KNOWN\_FAIL exclusion policy (9 specs), and the L0-  
1674 conditional augmentation campaign design.

### 1675 7. Error bars

1676 Question: Does the paper report error bars suitably and correctly?

1677 Answer: [Yes]

1678 Justification: Per-record pass rates include Wilson 95% confidence intervals; task-level  
1679 pass@ $k$  uses normal-approximation CIs on macro-averaged per-task probabilities. Direction  
1680 asymmetry uses McNemar’s exact test with Bonferroni correction (family-wise  $\alpha = 0.05$ ,  
1681 per-comparison  $\alpha = 0.01$  for five direction pairs). Effect sizes reported via Cohen’s  $h$ .  
1682 Cross-model comparison uses omnibus chi-squared, pairwise Fisher’s exact with Bonferroni  
1683 correction, McNemar’s test for paired concordance, and Cohen’s  $h$  effect sizes.

## 1684 8. Compute

1685 Question: For each experiment, does the paper provide sufficient information on the com-  
1686 puter resources used?

1687 Answer: [Yes]

1688 Justification: Table 19 specifies the evaluation platform (RTX 4070, Ryzen 9 7900X, Ubuntu  
1689 24.04, HPC SDK 24.3). Appendix G reports token usage and wall-clock time for all three  
1690 model campaigns.

## 1691 9. NeurIPS code of ethics

1692 Question: Does the research conform to the NeurIPS Code of Ethics?

1693 Answer: [Yes]

1694 Justification: The work evaluates publicly available benchmark code using commercial LLM  
1695 APIs. No human subjects or private data are involved. No immediate high-risk dual-use  
1696 concern beyond general automated coding capabilities; all benchmarks are public HPC  
1697 kernels under open-source licenses.

## 1698 10. Broader impacts

1699 Question: Does the paper discuss both potential positive societal impacts and negative  
1700 societal impacts of the work?

1701 Answer: [Yes]

1702 Justification: Section 6 discusses implications for HPC code migration and identifies the  
1703 limitation that LLM-generated parallel code should not be trusted without verification. The  
1704 benchmark itself enables systematic measurement of translation quality, which supports  
1705 responsible deployment.

## 1706 11. Safeguards

1707 Question: Does the paper describe safeguards that have been put in place for responsible  
1708 release of data or models?

1709 Answer: [N/A]

1710 Justification: PARBENCH releases benchmark specifications and evaluation infrastructure,  
1711 not trained models or generated datasets. The benchmark code originates from established  
1712 open-source HPC suites.

## 1713 12. Licenses

1714 Question: Are the creators of assets used in the paper properly credited and are the license  
1715 and terms of use explicitly mentioned and properly respected?

1716 Answer: [Yes]

1717 Justification: All five benchmark suites are cited (Rodinia Che et al. [2009], XSBench,  
1718 RSBench, mixbench, HeCBench). Repository URLs and commit hashes are recorded in  
1719 kernel specifications. LLM providers (Together AI, Azure OpenAI) are identified. Ro-  
1720 dinia is distributed under the Rodinia license (BSD-like); XSBench and RSBench under  
1721 MIT; HeCBench under BSD-3-Clause; mixbench under GPL-2.0. All are compatible with  
1722 academic use and redistribution.

## 1723 13. New assets

1724 Question: Are new assets introduced in the paper well documented and is the documentation  
1725 provided alongside the assets?

1726 Answer: [Yes]

1727 Justification: PARBENCH’s specification schema is documented (Section 2.1, Appendix B.1).  
 1728 The augmentation engine, evaluation pipeline, and harness are described in Section 2. The  
 1729 submitted artifact contents, hosting plan, and license summary are provided in Appendix I.

1730 **14. Crowdsourcing and human subjects**

1731 Question: For crowdsourcing experiments and research with human subjects, does the paper  
 1732 include the full text of instructions given to participants?

1733 Answer: [N/A]

1734 Justification: No crowdsourcing or human subjects research.

1735 **15. IRB approvals**

1736 Question: Does the paper describe potential risks incurred by study participants?

1737 Answer: [N/A]

1738 Justification: No human subjects research.

1739 **16. LLM usage declaration**

1740 Question: Does the paper disclose the use of LLMs in the core methodology?

1741 Answer: [Yes]

1742 Justification: Three LLMs (Qwen 3.5 397B-A17B, GPT-5.4, and GPT-5.3-codex) are the  
 1743 primary subjects of evaluation. Model identifiers, providers, and configurations are fully  
 1744 disclosed in Table 18 and Section 4.